
Limelight Documentation

Release 1.0

Limelight

Jan 12, 2019

1	Getting Started	2
1.1	Mounting	2
1.2	Wiring	4
1.3	Imaging	5
1.4	Networking Setup	6
1.5	Basic Programming	7
2	Vision Pipeline Tuning	10
2.1	Input	10
2.2	Thresholding	11
2.3	Contour Filtering	12
2.4	Output	13
3	Crosshair Calibration	14
3.1	Single Crosshair Mode:	14
3.2	Dual Crosshair Mode	14
4	Using Multiple Pipelines	15
5	Additional Theory	16
5.1	Vision Targets	16
5.2	Thresholding	17
5.3	Contour Filtering	19
5.4	From Pixels to Angles	21
6	Complete NetworkTables API	24
6.1	Advanced Usage with Raw Contours	26
7	Case Study: Estimating Distance	27
7.1	Using a Fixed Angle Camera	27
7.2	Using Area to Estimate Distance	28
8	Case Study: Aiming Using Vision	29
9	Case Study: Seeking	32
10	Case Study: Getting in Range	34

11 Case Study: Aiming and Range at the same time.	36
12 Case Study: 2017 Fuel Robot	38
13 Using Grip to create Custom OpenCV Vision Pipelines	44
13.1 Accessing the Video Stream within GRIP	45
13.2 How Limelight Handles Inputs to the Pipeline	46
13.3 Generating output for the Limelight to use	46
13.4 Code Generation for Limelight	47
14 Running GRIP Pipelines on Limelight	48
15 Frequently Asked Questions	51
15.1 Why is limelight using a low (320x240) resolution?	51
15.2 What if the game calls for a different tracking algorithm?	51
15.3 Why is there an extra usb port?	51
15.4 How do I view the video stream?	51
15.5 Are the eight LEDs bright enough?	52
15.6 Does Limelight support protocols other than NetworkTables?	52
15.7 Does Limelight support Active PoE?	52
15.8 Should I plug Limelight into our VRM (Voltage Regulator Module)?	52
15.9 My robot has extreme voltage swings while driving, will this damage my Limelight?	52
15.10 Will Limelight's LEDs dim when our robot's voltage drops?	52
16 Additional Resources	53
16.1 Team 254 Vision Talk	53
16.2 WPI Vision Processing Resources	53
16.3 GRIP	53
17 Software Change Log	54
17.1 2019.2 (1/7/19)	54
17.2 2019.1 (12/18/18)	55
17.3 2019.0 (12/10/18)	55
17.4 2018.5 (3/28/18)	55
17.5 2018.4 (3/19/18)	56
17.6 2018.3 (2/28/18)	56
17.7 2018.2 (2/10/18)	56
17.8 2018.1 (1/8/18)	58
17.9 2018.0 (1/3/18)	58
17.10 2017.7 (11/21/17)	60
17.11 2017.6 (11/13/17)	60
17.12 2017.5 (11/9/17)	61
17.13 2017.4 (10/30/17)	61
17.14 2017.3 (10/25/17)	61
17.15 2017.2 (10/23/17)	61
17.16 2017.1 (10/21/17)	61

Welcome to the official Limelight documentation.

Limelight is an easy-to-use smart camera for FRC. Refer to the table of contents below to get started.

CHAPTER 1

Getting Started

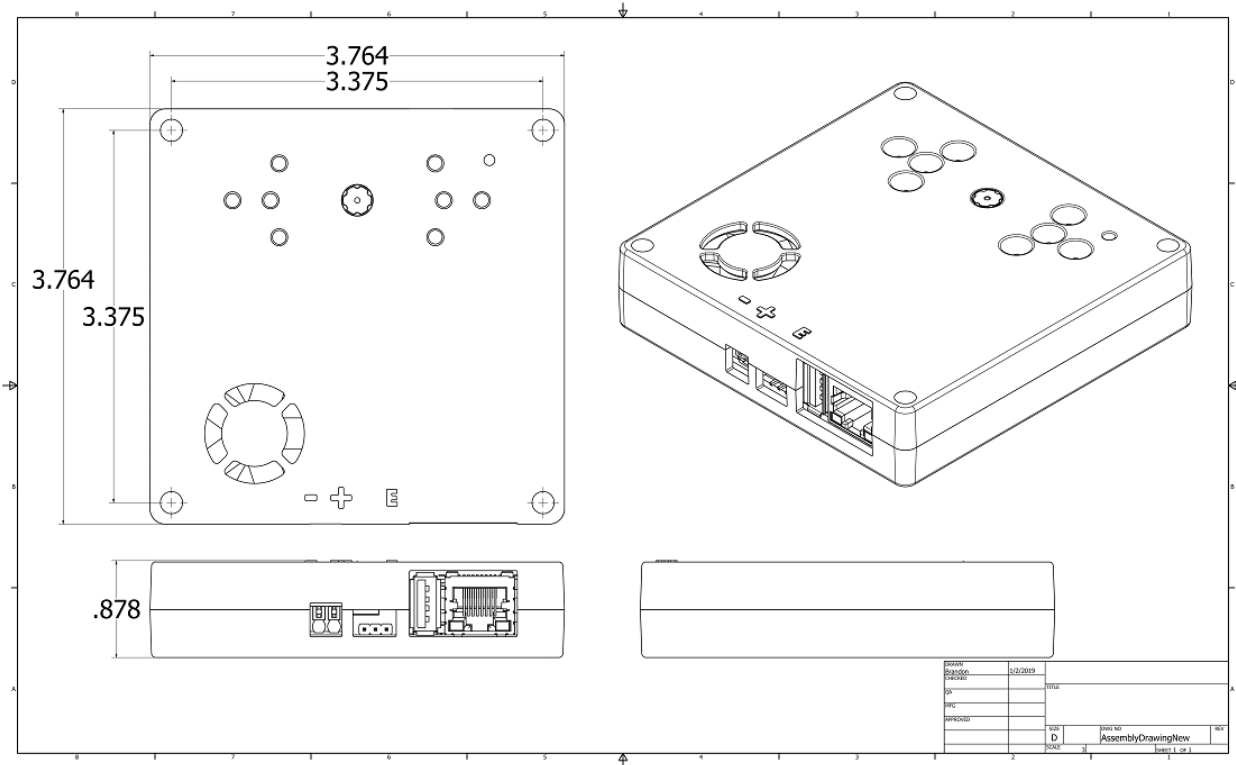
This page will cover the content on the official website's "Getting Started" page.

- *Mounting*
- *Wiring*
- *Imaging*
- *Networking Setup*
- *Basic Programming*

1.1 Mounting

Limelight 2

Use four 1 1/2" 10-32 screws and nylock nuts to mount your Limelight.



1.2 Wiring

Note: Limelight takes a 12V input, but is built to function down to 6V. Its LEDs have a constant brightness down to 7V.

Standard Wiring

- Do not run wires to your VRM.
- Run two wires from your limelight to a slot on your PDP (NOT your VRM).
- Add any breaker (5A, 10A, 20A, etc.) to the same slot on your PDP.
- Run an ethernet cable from your Limelight to your robot radio.

Power-over-Ethernet (PoE) Wiring

Note: PoE allows you to add both power and network connectivity to your Limelight via an Ethernet cable.

Warning: This is not standard 44V PoE - this is why you must use a passive injector with 12V.

- (LIMELIGHT 1 ONLY) Ensure that your Limelight's power jumper is set to the "E" position.
- Connect a passive [Passive PoE Injector](#) to your PDP (NOT your VRM).

- Add any breaker (5A, 10A, 20A, etc.) to the same slot on your PDP.
- Run an ethernet cable from your Limelight to your passive POE injector.

1.3 Imaging

Note: Limelight will receive software updates for improvements and features to accomodate the game if necessary.

Limelight 2

- Do not use a Windows 7 machine.
- Remove power from your limelight.
- Download the latest drivers, flasher tool, and image from from the [Downloads](#) Page.
- Install the Balena Etcher flash tool.
- Run a USB-MicroUSB cable from your laptop to your limelight.
- Run “Balena Etcher”.
- It may take up to 20 seconds for your machine to recognize the camera.
- Select the latest .zip image in your downloads folder
- Select an “RPI” device in the “Limelights” menu
- Click “Flash”
- Once flashing is complete, remove power from your limelight

Limelight 1

- Do not use a Windows 7 machine.
- Remove power from your limelight.
- Download the latest drivers, flasher tool and image from from the [Downloads](#) Page.
- Install the Balena Etcher flash tool.
- Run a USB-MicroUSB cable from your laptop to your limelight.
- Apply power to your limelight.
- Run “Balena Etcher”.
- It may take up to 20 seconds for your machine to recognize the camera.
- Select the latest .zip image in your downloads folder
- Select an “RPI” device in the “Limelights” menu
- Click “Flash”
- Once flashing is complete, remove power from your limelight

Warning: Only connect the microUSB cable while imaging. Limelight enters a special flash mode while the microUSB cable is connected. You will not be able to access the web interface while Limelight is in flash mode.

1.4 Networking Setup

We highly recommend following the static IP instructions for reliability during events.

Follow these steps before starting:

- Go to add/remove programs in windows, and search for “bonjour”
- **How many items do you see?**
 - If there are two (2) items with “bonjour” in their name, uninstall “bonjour print services”
 - If there are no (0) items with “bonjour” in their name, install Bonjour from our Downloads page.
- Reboot your robot and computer.
- Follow the steps listed below.

Static IP Address (Recommended)

- Install Bonjour from the [Downloads Page](#).
- Power-up your robot, and connect your laptop to your robot’s network.
- After your Limelight flashes its LED array, navigate to <http://limelight.local:5801>. This is the configuration panel.
- Navigate to the “Settings” tab on the left side of the interface.
- Enter your team number and press the “Update Team Number” button.
- Change your “IP Assignment” to “Static”.
- Set your Limelight’s IP address to “10.TE.AM.11”.
- Set the Netmask to “255.255.255.0”.
- Set the Gateway to “10.TE.AM.1”.
- Click the “Update” button.
- Power-cycle your robot.
- You will now be able to access your config panel at <http://10.TE.AM.11:5801>, and your camera stream at <http://10.TE.AM.11:5800>

Dynamic IP Address (Not recommended)

- Install Bonjour from the [Downloads Page](#).
- Power-up your robot, and connect your laptop to your robot’s network.
- After your Limelight flashes its LED array, navigate to <http://limelight.local:5801>. This is the configuration panel.
- Navigate to the “Settings” tab on the left side of the interface.
- Enter your team number and press the “Update Team Number” button.
- Change your “IP Assignment” to “Automatic”.
- Click the “Update” button.
- Power-cycle your robot.
- You can continue to access your config panel at <http://limelight.local:5801>, and your camera stream at <http://limelight.local:5800>

Note:

17. Why do we recommend a static IP?

1. First, it shaves multiple seconds off Limelight's boot time. Second, teams have historically had issues with DHCP assignment and mDNS responders on actual FRC fields and with event radio firmware.

Note:

17. How do I reset the IP address?

1. After your Limelight has booted, hold the config button on the front face of the camera until the LEDs start blinking. Power-cycle your robot, and your Limelight will have an automatically-assigned IP address.

- If the above steps do not fix the problem, install Angry IP scanner and find the address for your limelight.
- Go to <limelightaddress>:5801, and give your limelight a .11 static IP.
- From this point onward, you can rely on the static IP to access the page.

1.5 Basic Programming

For now, we just need to get data from the camera to your robot. Limelight posts targeting data to Network Tables at 100hz. The default update rate for NetworkTables is 10hz, but Limelight automatically overwrites it to allow for more frequent data transfer.

To get started, we recommend reading four values from the "limelight" Network Table at at least 100hz. Our code samples will show you exactly how to do this. The offsets to your target (in degrees) are sent as "tx" and "ty". These can be used to turn your robot, turn a turret, etc. The target's area, sent as "ta", may be used a rough indicator of distance to your target. Area is a value between 0 and 100, where 0 means that your target's hull area is 0% of the total image area, and 100 means that your target's hull fills the entire image. The rotation or "skew" of your target is returned as "ts". If all values are equal to zero, no targets exist.

In addition, you may control certain features by setting values in NetworkTables.

Read the following from the "limelight" table

tv	Whether the limelight has any valid targets (0 or 1)
tx	Horizontal Offset From Crosshair To Target (-27 degrees to 27 degrees)
ty	Vertical Offset From Crosshair To Target (-20.5 degrees to 20.5 degrees)
ta	Target Area (0% of image to 100% of image)
ts	Skew or rotation (-90 degrees to 0 degrees)
tl	The pipeline's latency contribution (ms) Add at least 11ms for image capture latency.

Write the following to the "limelight" table

ledMode	Sets limelight's LED state
0	use the LED Mode set in the current pipeline
1	force off
2	force blink
3	force on

camMode	Sets limelight's operation mode
0	Vision processor
1	Driver Camera (Increases exposure, disables vision processing)

pipeline	Sets limelight's current pipeline
0 .. 9	Select pipeline 0..9

Java

```
NetworkTable table = NetworkTableInstance.getDefault().getTable("limelight");
NetworkTableEntry tx = table.getEntry("tx");
NetworkTableEntry ty = table.getEntry("ty");
NetworkTableEntry ta = table.getEntry("ta");

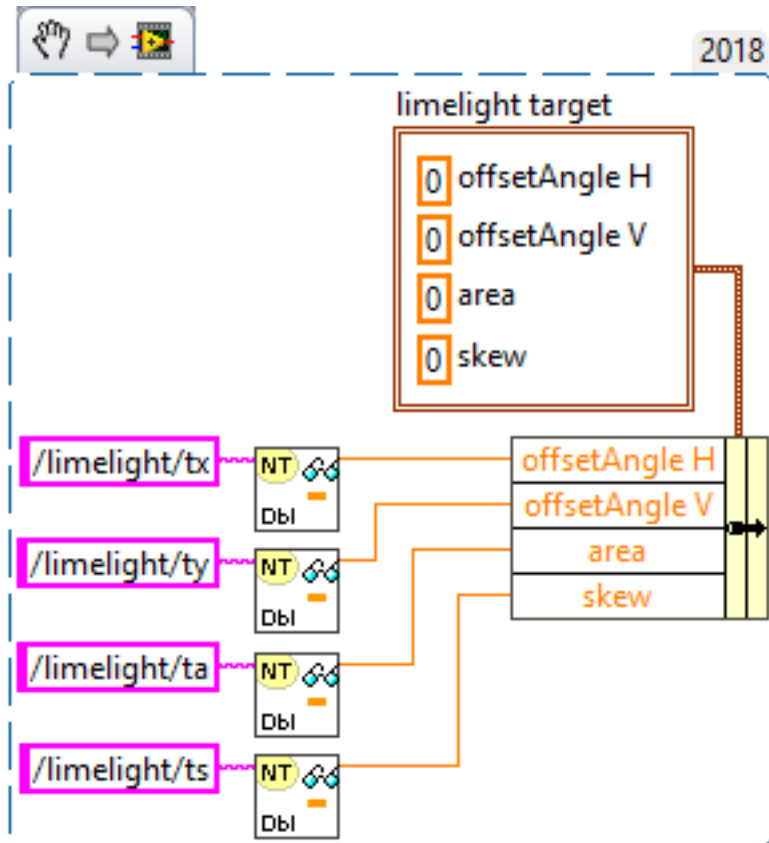
//read values periodically
double x = tx.getDouble(0.0);
double y = ty.getDouble(0.0);
double area = ta.getDouble(0.0);

//post to smart dashboard periodically
SmartDashboard.putNumber("LimelightX", x);
SmartDashboard.putNumber("LimelightY", y);
SmartDashboard.putNumber("LimelightArea", area);
```

Don't forget to add these imports:

```
import edu.wpi.first.wpilibj.smartdashboard.SmartDashboard;
import edu.wpi.first.networktables.NetworkTable;
import edu.wpi.first.networktables.NetworkTableEntry;
import edu.wpi.first.networktables.NetworkTableInstance;
```

LabView



C++

```
std::shared_ptr<NetworkTable> table = nt::NetworkTableInstance::GetDefault().GetTable(
    "limelight");
double targetOffsetAngle_Horizontal = table->GetNumber("tx", 0.0);
double targetOffsetAngle_Vertical = table->GetNumber("ty", 0.0);
double targetArea = table->GetNumber("ta", 0.0);
double targetSkew = table->GetNumber("ts", 0.0);
```

Don't forget to add these includes:

```
include "frc/smartdashboard/Smartdashboard.h"
include "networktables/NetworkTable.h"
```

Python

```
from networktables import NetworkTables

table = NetworkTables.getTable("limelight")
tx = table.getNumber('tx', None)
ty = table.getNumber('ty', None)
ta = table.getNumber('ta', None)
ts = table.getNumber('ts', None)
```

Vision Pipeline Tuning

To configure the limelight vision pipeline, first access its web interface at <http://10.te.am.11:5801>. If you have opted to use a dynamically assigned ip-address, access the interface at <http://limelight.local:5801>.

The “Tracking” page is comprised of four tuning tabs:

- *Input*
- *Thresholding*
- *Contour Filtering*
- *Output*

In addition, you can adjust your Limelight’s crosshair by pressing the “calibrate crosshair” button while a target is in view. In a future release, crosshair calibration will move to the “output” tab.

2.1 Input

The Input Tab hosts controls to change the raw camera image before it is passed through the processing pipeline.

2.1.1 Orientation

Controls the orientation of incoming frames. Set it to “inverted” if your camera is mounted upside-down.

2.1.2 Exposure

Controls the camera's exposure setting in milliseconds. Think of a camera as a grid of light-collecting buckets - exposure time controls how long your camera's "buckets" are open per frame. Lowering the exposure time will effectively darken your image. Low and fixed exposure times are crucial in FRC, as they black-out the bulk of incoming image data. Well-lit retroreflective tape will stand out in a mostly black image, making vision processing a straightforward process.

2.2 Thresholding

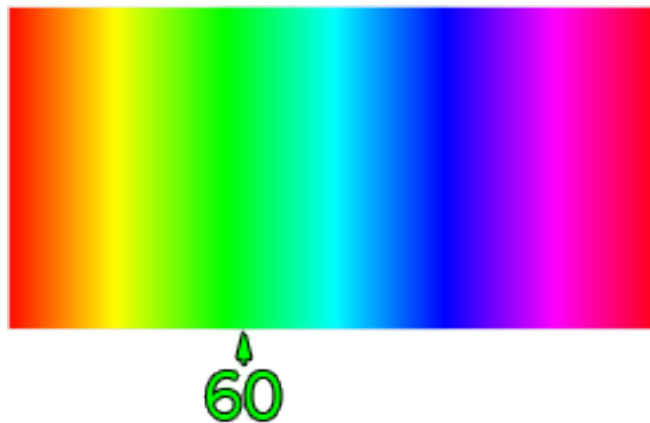
Thresholding is a critical component of most FRC vision tracking algorithms. It is the act of taking an image, and throwing away any pixels that aren't in a specific color range. The result of thresholding is generally a one-dimensional image in which a pixel is either "on" or "off."

2.2.1 Video Feed

Controls which image is streamed from the mjpeg server. You should switch to the "threshold" image if you need to tune your HSV thresholding.

2.2.2 Hue

Describes a "pure" color. A Hue of "0" describes pure red, and a hue of 1/3 (59 on the slider) describes pure green. Hue is useful because it doesn't change as a pixel "brightens" or "darkens". This is the most important parameter to tune. If you make your hue range as small as possible, you will have little if any trouble transitioning to an actual FRC field.



2.2.3 Saturation

Describes the extent to which a color is "pure". Another way to think of this is how washed-out a color appears, that is, how much "white" is in a color. Low saturation means a color is almost white, and high saturation means a color is almost "pure".

2.2.4 Value

Describes the darkness of a color, or how much “black” is in a color. A low value corresponds to a near-black color. You should absolutely increase the minimum value from zero, so that black pixels are not passed through the processing pipeline.

2.2.5 Thresholding Wands

Wands enable users to click on Limelight’s video stream to perform automatic HSV thresholding.

- The “Set” wand centers HSV parameters around the selected pixel
- The “Add” wand adjusts HSV parameters to include the selected pixel
- The “Subtract” wand adjust HSV paramters to ignore the selected pixel

2.2.6 Erosion and Dilation

Erosion slightly erodes the result of an HSV threshold. This is useful if many objects are passing through a tuned HSV threshold. Dilation slightly inflates the result of an HSV threshold. Use this to patch holes in thresholding results.

*GIF

2.3 Contour Filtering

After thresholding, Limelight applies a to generate a list of contours. After that, each contour is wrapped in a bounding rectangle, or “convex hull”. Hulls are passed through a series of filters to determine the “best” hull. If multiple hulls pass through all filters, Limelight chooses the largest hulls.

2.3.1 Target Area

Controls the range of acceptable bounding-rectangle areas, as percentages of the screen. You can increase the minimum area to help filter-out stadium lights, and decrease the maximum value to help filter-out things like large displays near the field.

Note: The area slider is not linearly scaled, but quarticly scaled. This is done to provide extra precision near the lower-end of area values, where many FRC targets lie. In any case, an area slider should be quadratically scaled, as the area of a square scales quadratically with its side length.

2.3.2 Target Fullness

Fullness is the percentage of “on” pixels in the chosen contour’s bounding rectangle. A solid rectangle target will have a near-1.0 fullness, while a U-shaped target will have a low fullness.

2.3.3 Target Aspect Ratio

Aspect ratio is defined by the width of the bounding rectangle of the chosen contour divided by its height. A low aspect ratio describes a “tall” rectangle, while a high aspect ratio describes a “wide” rectangle.

Note: The aspect ratio slider is also quadratically scaled.

2.4 Output

This tab controls what happens during the last stage of the vision pipeline

2.4.1 Targeting Region

Controls the point of interest of the chosen contour’s bounding rectangle. By default, the tracking parameters tx and ty represent the offsets from your crosshair to the center of the chosen rectangle. You can use another option if a target changes in size, or is comprised of two targets that sometimes blend together.

2.4.2 Target Grouping for Combined Targets

Controls target “grouping”. Set to dual mode to look for “targets” that consist of two shapes, or tri mode to look for targets that consist of three shapes.

2.4.3 Crosshair Calibration

Controls the “origin” of your targeting values. Let’s say a shooter on your robot needs to be calibrated such that it always points a bit left-of-center. You can line up your robot, click “calibrate,” and all of your targeting values will be sent relative to your new crosshair. See the calibration page for more details!

Crosshair Calibration

- *Single Crosshair Mode:*
- *Dual Crosshair Mode*

Limelight's crosshairs turn calibration into a seamless process. Rather than storing offsets in their code, teams can line-up their robots perfectly by hand (or by joystick), and click the "calibrate" button.

Calibrating a crosshair moves the "zero" of your targeting data. This is very useful if your Limelight isn't perfectly centered on your robot.

3.1 Single Crosshair Mode:

Line-up your robot at its ideal scoring location+rotation, and click "calibrate". Now a tx and ty of "zero" equate to a perfectly aligned robot. If your robot needs to be recalibrated for a new field, simply take a practice match to find the perfect alignment for your robot, and click "calibrate" during your match.

3.2 Dual Crosshair Mode

Imagine a robot with an off-axis camera or shooter that needs to shoot gameobjects into a goal from many positions on the field. As the robot approaches the goal, its crosshair must adjust in real-time to compensate. Dual crosshair mode is built for this functionality. Line-up your robot at its closest scoring position+rotation, and calibrate crosshair "A". Line-up your robot at its farthest scoring position+rotation, and calibrate corsshair "B". When you calibrate in dual-crosshair mode, crosshairs also store an area value. You will notice that as your robot moves between its min and max scoring distances, the crosshair moves between crosshair "A" and crosshair "B". This is done by checking the area of the target, and comparing it to the two target areas seen during calibration.

Using Multiple Pipelines

Limelight can store up to ten unique vision pipelines for different goals, different fields, or different robots. Change pipelines mid-match by changing the “pipeline” value in NetworkTables.

To edit multiple pipelines, you must first check the “Ignore NetworkTables Index” checkbox in the web interface. This will force the robot to temporarily allow you to change the pipeline index through the webinterface rather than through NetworkTables.

To download your pipelines for backups and sharing, simply click the “download” button next to your pipeline’s name. To upload a pipeline, click the “upload” button.

Here’s an example of a robot that utilizes two pipelines:

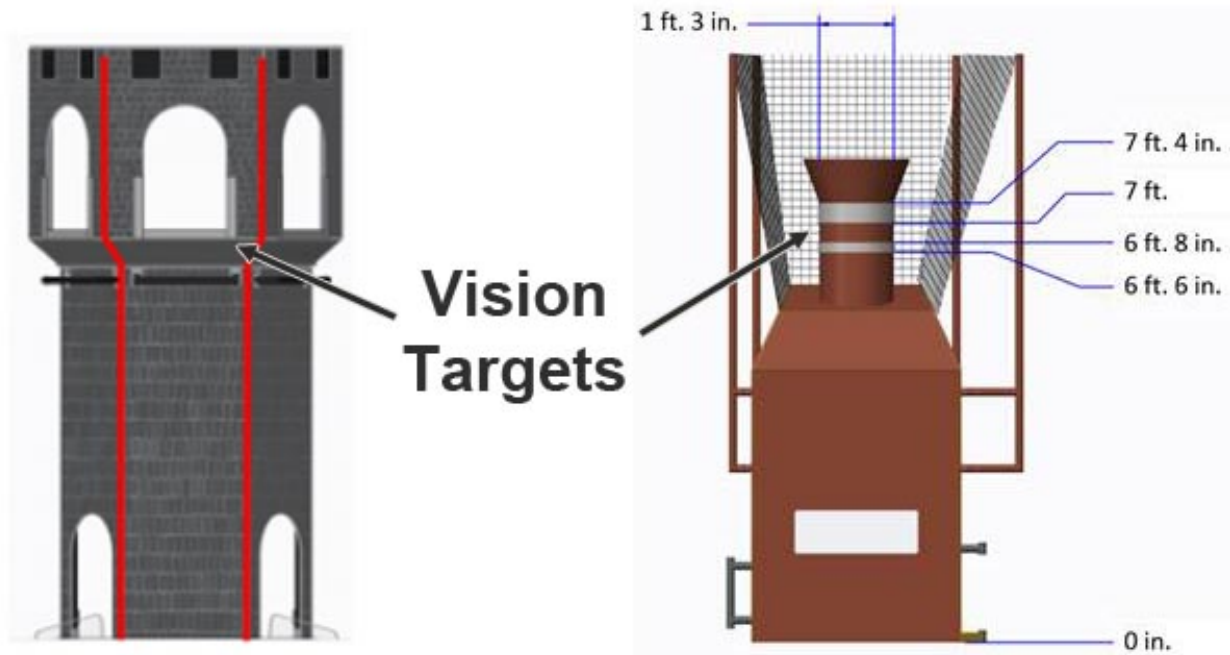
The first pipeline is tuned to target single vertical stripes. The second pipeline is tuned to find a combo of two horizontal stripes. The code for this robot is available in the “Aim and Range” case study.

Notice that when the robot switches pipelines, the web interface auto-loads the new pipeline.

- Vision-Targets
- *Thresholding*
- *Contour Filtering*
- *Output*

5.1 Vision Targets

The FRC game designers often place reflective “vision targets” on the field in strategic locations. These vision targets are usually made out of retro-reflective tape. Major scoring elements generally have vision targets that can be used to automatically aim. Below you can see two examples of some of the vision targets from the 2016 and 2017 FRC games.



These retro-reflective vision targets have a very useful property: when light is shined at them, it will reflect directly back to the light source. This is why Limelight has bright green LEDs surrounding its camera lens. By setting the camera exposure very low while emitting a bright green light toward the target, we can acquire an image that is mostly black with a bright green vision target. This makes the job of acquiring the target relatively easy.

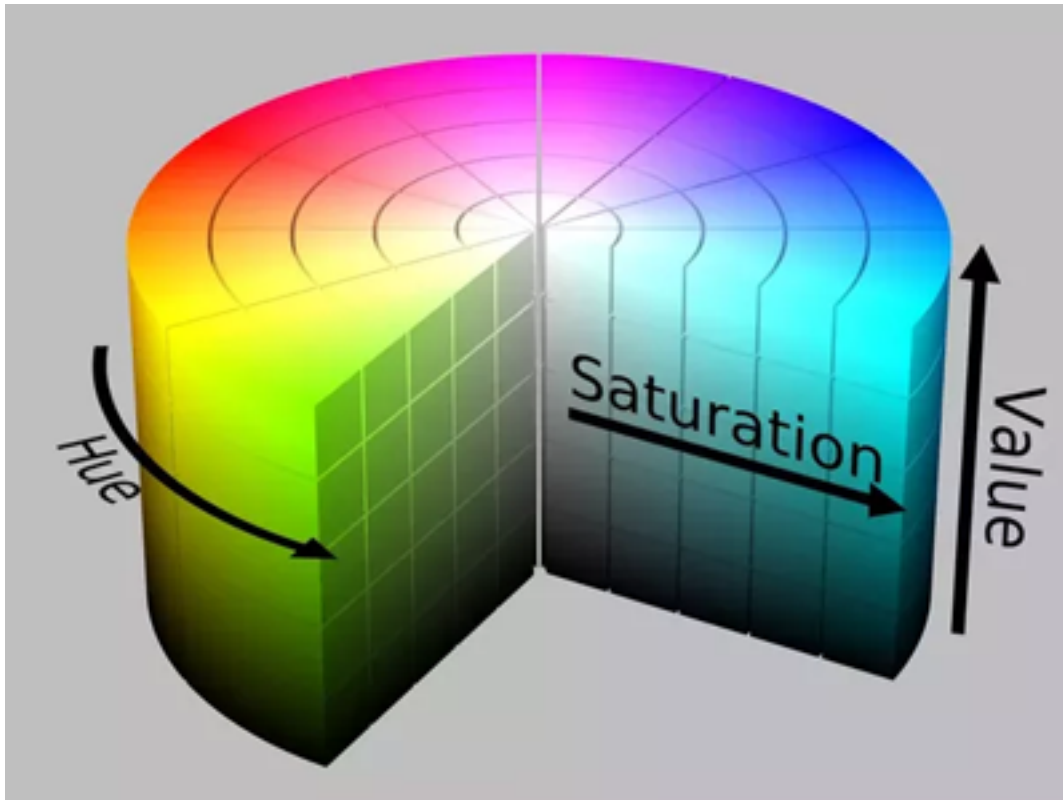
Here you can see an example of the ideal of image. Notice how almost all of the detail in the image is gone due to the low exposure setting but the retro-reflective tape stands out brightly.



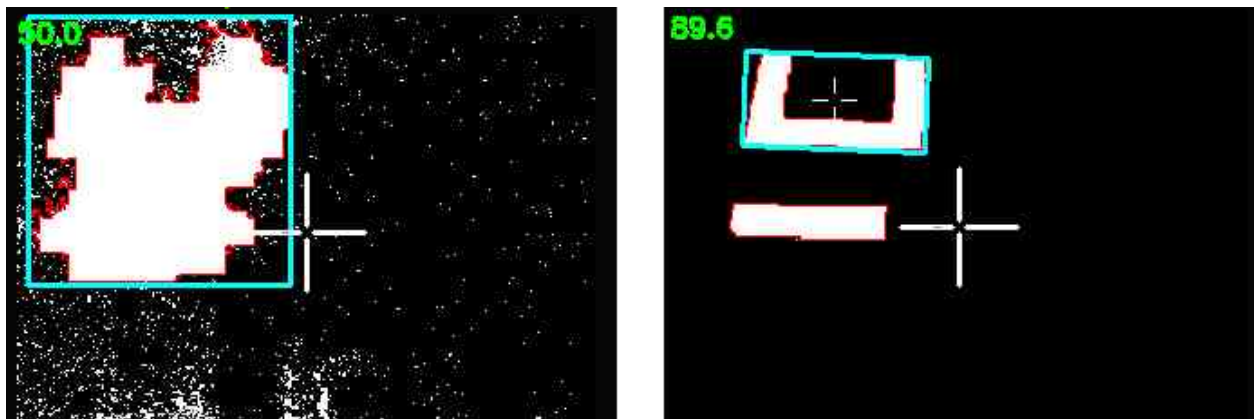
5.2 Thresholding

Thresholding is the next critical component of most FRC vision tracking algorithms. It is the act of taking an image, and throwing away any pixels that aren't in a specific color range. The result of thresholding is generally a one-dimensional image in which a pixel is either "on" or "off". Thresholding works very well on images that are captured using the above strategy (low exposure, very dark image with a brightly illuminated vision target)

Limelight does thresholding in the HSV (Hue-Saturation-Value) colorspace. You may be used to thinking of colors in the RGB (Red-Green-Blue) colorspace. HSV is just another way of representing color similar to the way cartesian coordinates or polar coordinates can be used to describe positions. The reason we use the HSV colorspace is that the Hue can be used to very tightly select the green color that Limelight's leds output.



It is critical to adjust your thresholding settings to eliminate as much as you can from the image. You will get the best results if you optimize each stage of your vision pipeline before moving to the next stage. The following image depicts the difference between improper and proper thresholding:



Sometimes things like ceiling lights or windows in an arena can be difficult to remove from the image using thresholding, which brings us to the next stage.

5.3 Contour Filtering

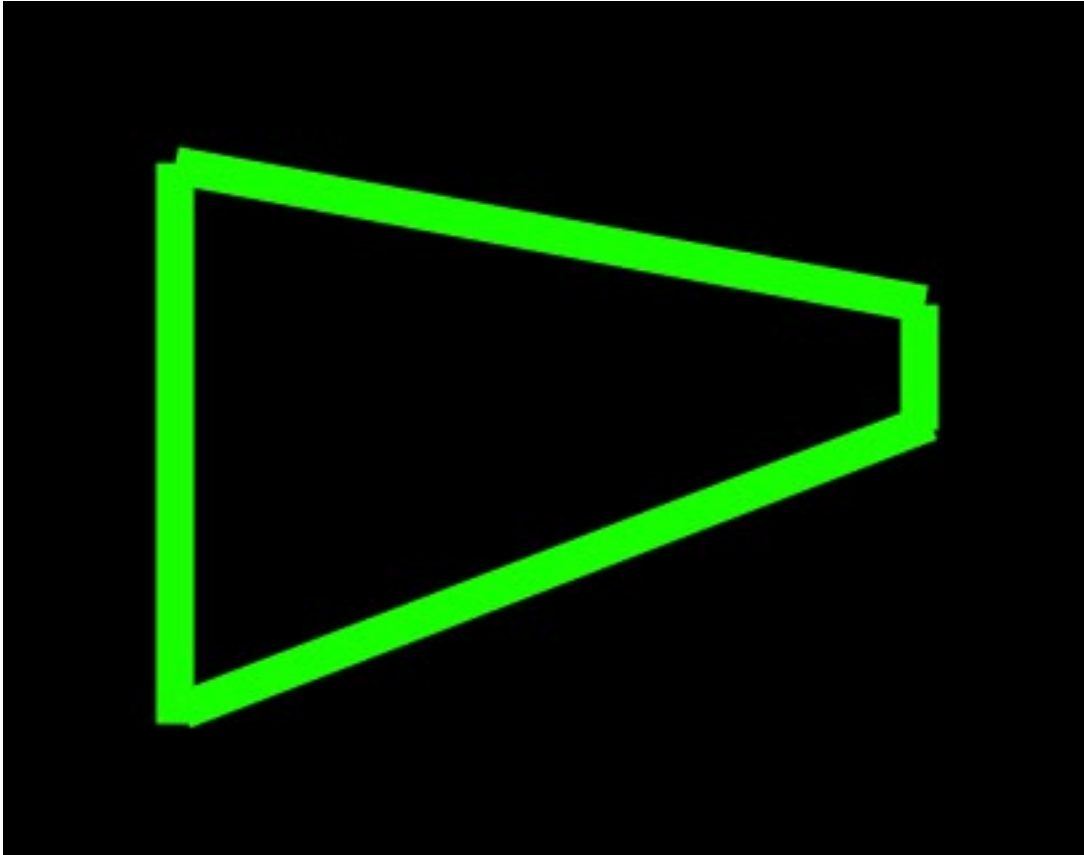
After thresholding, Limelight's vision pipeline generates a set of contours for the image. A contour is a curve surrounding a contiguous set of pixels. Sometimes things like ceiling lights, arena scoreboards, windows and other things can make it past the thresholding step. This is where contour filtering becomes useful. The goal is to eliminate any contours which we know are not the target we are interested in.

The first and easiest contour filter is to ignore any contours which are smaller than what our vision target looks like from our scoring distance. Anything smaller than that size is obviously something farther away and should be ignored. This is called area filtering.

The FRC vision targets often have some geometric property that can be exploited to help us filter contours. For example, if the vision target has a wide aspect ratio, we can filter out any contours that are not wide:



However, keep in mind that your camera may be looking at the target from an odd angle. This can drastically affect the aspect ratio of its contour. Be sure to test your settings from a variety of angles to ensure that you do not filter too aggressively and end up ignoring the vision target!



This next image target is very interesting. It is one of the best designed vision targets in FRC (in our opinion). Limelight automatically calculates a value called the **fullness** of a contour. **Fullness** is the ratio between the pixel area of the contour to its convex area. This particular shape has a very low fullness and you almost never see any ceiling lights, windows, etc with such a low fullness. So you can very effectively filter out the unwanted contours if your vision target looks like this one.



Limelight has many options for filtering contours. You can use these options along with what you know about the geometry properties of the particular vision target you are trying to track.

Currently, if multiple contours pass through your filtering options, the largest contour is chosen. In addition, the pipeline prefers to “lock” onto contours using hysteresis. This is a feature that helps prevent flickering between similar targets.

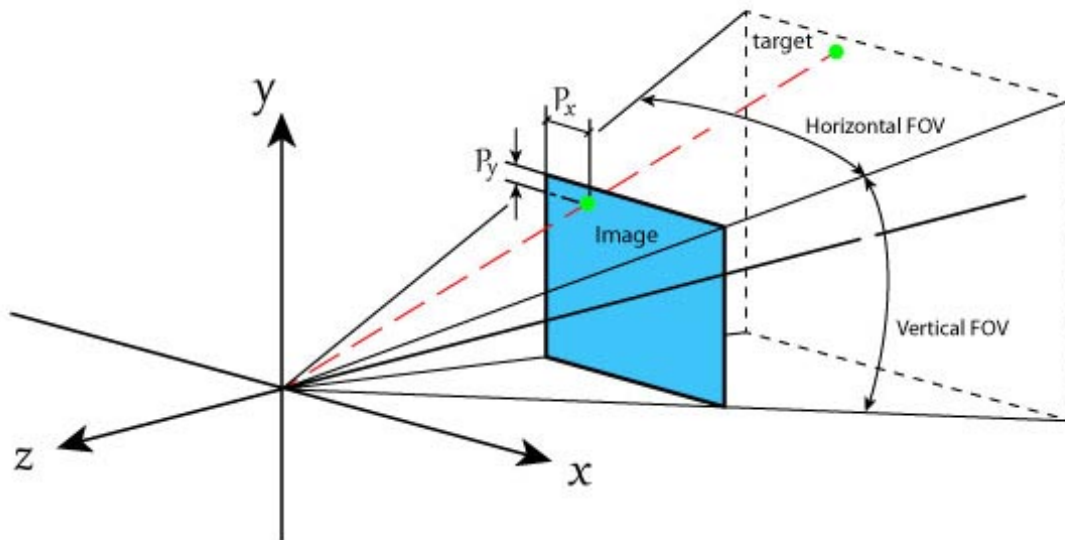
5.4 From Pixels to Angles

The end result of the vision pipeline is a pixel location of the best contour in the image. For most games, we can just aim at the center of the contour. Sometimes it is also useful to aim at the top-center or some other point but essentially we have a pixel coordinate for where we want to aim. In order to compute the angles to this target, we need to use a little bit of trigonometry.

First we assume that we have a perfect “pinhole” camera. In practice this can be far from the truth but the limelight’s camera is very close. (A fisheye lens would be far from this ideal as a counter-example.)

The limelight camera has a horizontal field of view of 54 degrees and a vertical field of view of 41 degrees. It captures images at 320x240 resolution. We assume the center of the image is the optical axis of the camera (so the x and y angles for that location are 0,0). Given these known values, we can use a little trigonometry to compute the angles for any pixel in the image.

The below diagram shows an example target point which we want to compute angles for. Pixel coordinates start at the upper left corner of the image and are positive to the right and down.



Our first step will be to convert from pixel coordinates to normalized 2D coordinates where 0,0 is the center of the image and 1.0:

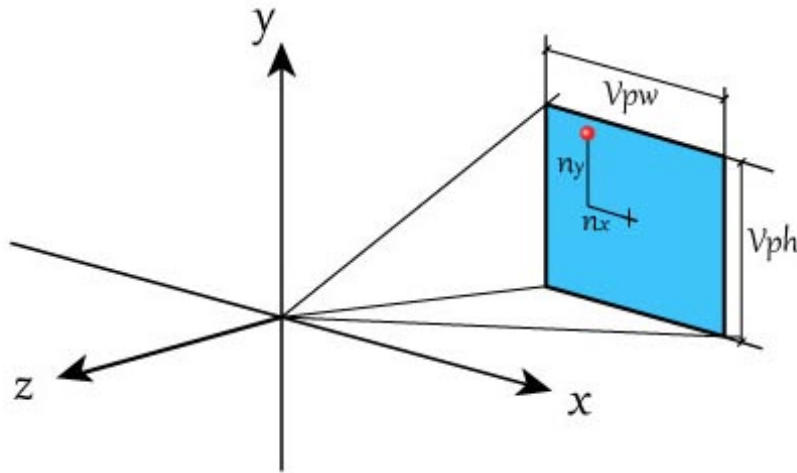
(px,py) = pixel coordinates, 0,0 is the upper-left, positive down and to the right

(nx,ny) = normalized pixel coordinates, 0,0 is the center, positive right and up

$$nx = (1/160) * (px - 159.5)$$

$$ny = (1/120) * (119.5 - py)$$

Next we define an imaginary view plane and compute its size. For simplicity, we can choose to place this plane 1.0 unit in front of the camera location. Here is a view looking down on the camera. Our goal is to compute the view plane width and height as those values will be used to compute the angles later:



Given a distance of 1.0 unit and a known horizontal and vertical fov, we can compute the size of the view plane rectangle the following formulas:

$$vpw = 2.0 * \tan(\text{horizontal_fov}/2)$$

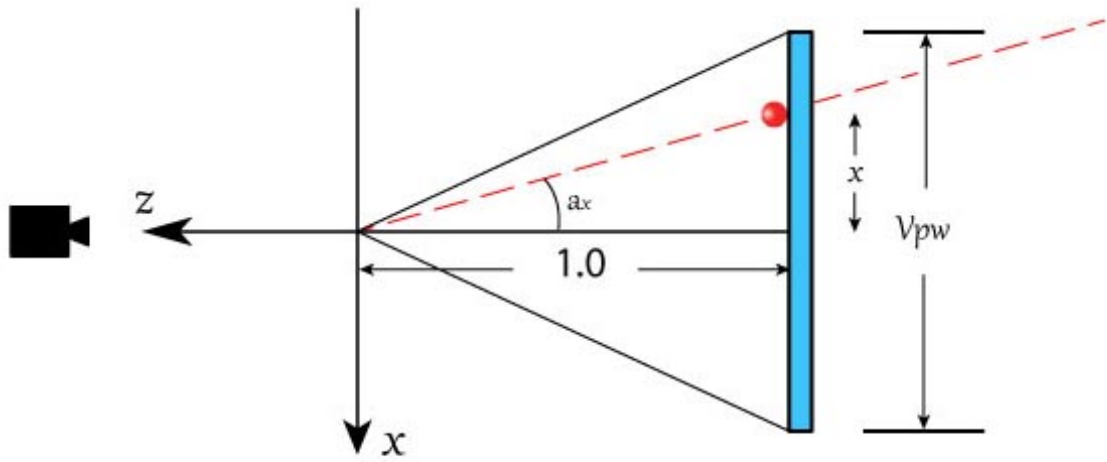
$$vph = 2.0 * \tan(\text{vertical_fov}/2)$$

Using these two values, we can now convert between normalized pixel coordinates and view plane coordinates using a simple multiply.

$$x = vpw/2 * nx;$$

$$y = vph/2 * ny;$$

Remember that we chose our view plane to be positioned at a distance of 1.0. Now we have everything we need to compute the angles to the target point.



$$\tan(ax) = x / 1$$

$$\tan(ay) = y / 1$$

$$ax = \text{atan2}(1, x)$$

$$ay = \text{atan2}(1, y)$$

CHAPTER 6

Complete NetworkTables API

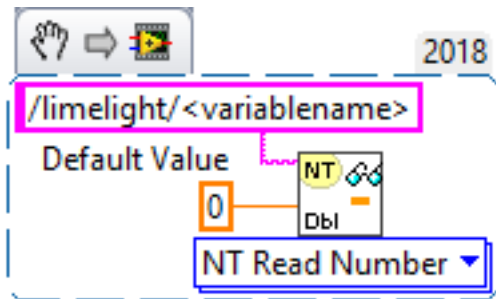
“Best” Contour information:

Use the following code:

Java

```
NetworkTableInstance.getDefault().getTable("limelight").getEntry("<variablename>").  
    ↪ getDouble(0);
```

LabView



C++

```
nt::NetworkTableInstance::GetDefault().GetTable("limelight")->GetNumber("<variablename>", 0.0);
```

Python

```
NetworkTables.getTable("limelight").getNumber('<variablename>');
```

to retrieve this data:

tv	Whether the limelight has any valid targets (0 or 1)
tx	Horizontal Offset From Crosshair To Target (-27 degrees to 27 degrees)
ty	Vertical Offset From Crosshair To Target (-20.5 degrees to 20.5 degrees)
ta	Target Area (0% of image to 100% of image)
ts	Skew or rotation (-90 degrees to 0 degrees)
tl	The pipeline's latency contribution (ms) Add at least 11ms for image capture latency.
tshort	Sidlength of shortest side of the fitted bounding box (pixels)
tlong	Sidlength of longest side of the fitted bounding box (pixels)
thoriz	Horizontal sidlength of the rough bounding box (0 - 320 pixels)
tvert	Vertical sidlength of the rough bounding box (0 - 320 pixels)

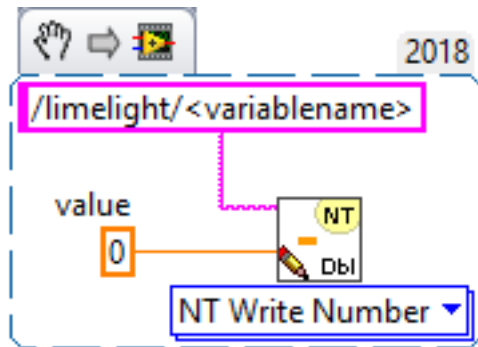
Camera Controls:

Use the following code:

Java

```
NetworkTableInstance.getDefault().getTable("limelight").getEntry("<variablename>").
    setNumber(<value>);
```

LabView



C++

```
nt::NetworkTableInstance::GetDefault().GetTable("limelight")->PutNumber("<variablename>", <value>);
```

Python

```
NetworkTables.getTable("limelight").putNumber('<variablename>', <value>);
```

to set this data:

ledMode	Sets limelight's LED state
0	use the LED Mode set in the current pipeline
1	force off
2	force blink
3	force on

camMode	Sets limelight's operation mode
0	Vision processor
1	Driver Camera (Increases exposure, disables vision processing)

pipeline	Sets limelight's current pipeline
0 .. 9	Select pipeline 0..9

stream	Sets limelight's streaming mode
0	Standard - Side-by-side streams if a webcam is attached to Limelight
1	PiP Main - The secondary camera stream is placed in the lower-right corner of the primary camera stream
2	PiP Secondary - The primary camera stream is placed in the lower-right corner of the secondary camera stream

snapshot	Allows users to take snapshots during a match
0	Stop taking snapshots
1	Take two snapshots per second

6.1 Advanced Usage with Raw Contours

Raw Targets:

Limelight posts three raw contours to NetworkTables that are not influenced by your grouping mode. That is, they are filtered with your pipeline parameters, but never grouped. X and Y are returned in normalized screen space (-1 to 1) rather than degrees.

tx0	Raw Screenspace X
ty0	Raw Screenspace Y
ta0	Area (0% of image to 100% of image)
ts0	Skew or rotation (-90 degrees to 0 degrees)
tx1	Raw Screenspace X
ty1	Raw Screenspace Y
ta1	Area (0% of image to 100% of image)
ts1	Skew or rotation (-90 degrees to 0 degrees)
tx2	Raw Screenspace X
ty2	Raw Screenspace Y
ta2	Area (0% of image to 100% of image)
ts2	Skew or rotation (-90 degrees to 0 degrees)

Raw Crosshairs:

If you are using raw targeting data, you can still utilize your calibrated crosshairs:

cx0	Crosshair A X in normalized screen space
cy0	Crosshair A Y in normalized screen space
cx1	Crosshair B X in normalized screen space
cy1	Crosshair B Y in normalized screen space

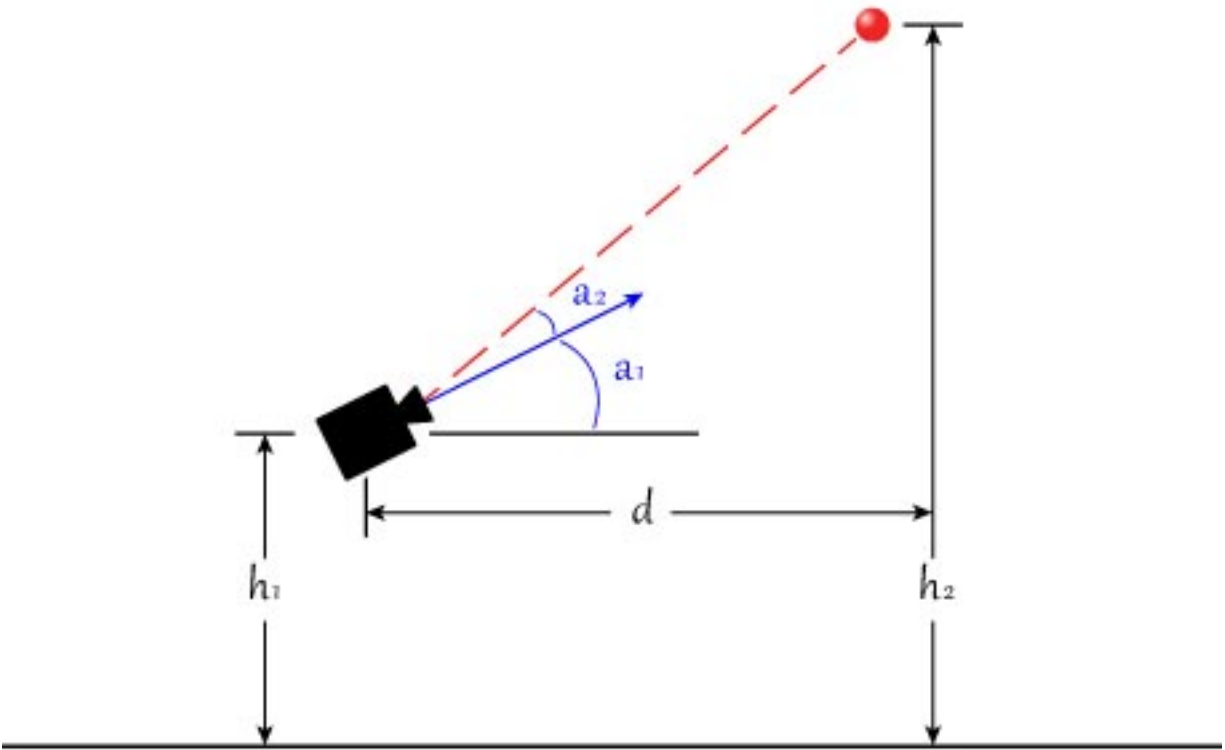
Case Study: Estimating Distance

- *Using a Fixed Angle Camera*
- *Using Area to Estimate Distance*

7.1 Using a Fixed Angle Camera

If your vision tracking camera is mounted on your robot in such a way that its angle between the ground plane and its line of sight does not change then you can use this technique to very accurately calculate the distance to a target. You can then use this distance value to either drive your robot forward and back to get into the perfect range or you can use it to adjust the power of a shooting mechanism.

See the below diagram. In this situation all of the variables are known: the height of the target (h_2) is known because it is a property of the field. The height of your camera above the floor (h_1) is known and its mounting angle is known (a_1). The limelight (or your vision system) can tell you the y angle to the target (a_2).



We can solve for d using the following equation:

$$\tan(a_1 + a_2) = (h_2 - h_1) / d$$

$$d = (h_2 - h_1) / \tan(a_1 + a_2)$$

When using this technique it is important to choose the mounting angle of your camera carefully. You want to be able to see the target both when you're too close and too far away. You also do not want this angle to change so mount it securely since you'll have to adjust all of your numbers if it moves.

If you are having trouble figuring out what the angle a_1 is, you can also use the above equation to solve for a_1 . Just put your robot at a known distance (measuring from the lens of your camera) and solve the same equation for a_1 .

In the case where your target is at nearly the same height as your camera, this technique is not useful.

7.2 Using Area to Estimate Distance

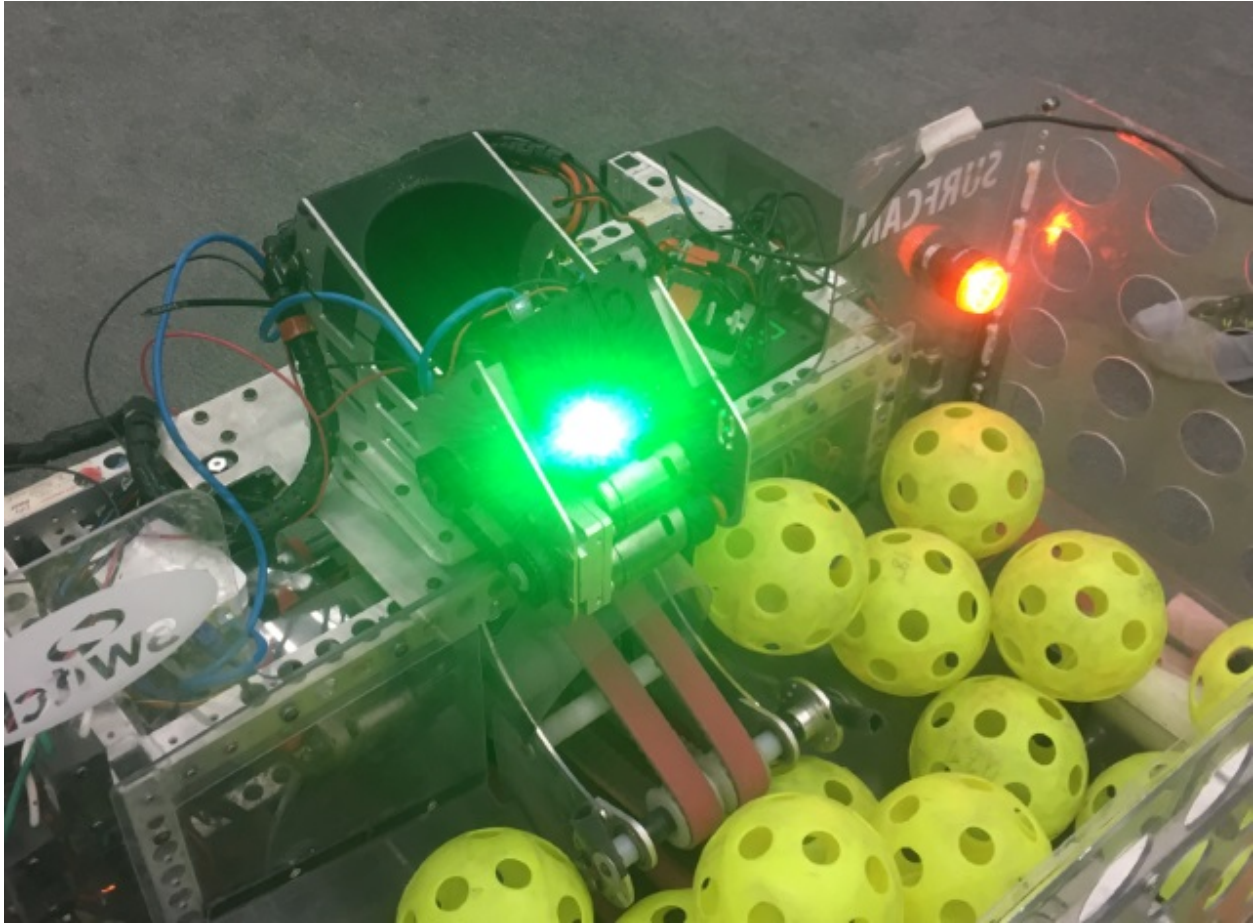
Another simple way to estimate distance is to use the area of the contour you are tracking. This is a very simple method to implement but it does not give you extremely accurate results. All you do is point your vision camera at the target from a known distance and take note of the area of the blob. Make sure you're using an accurate representation of the real field's vision target and make sure you're pointing at it from your desired shooting location. You can then do this from a few different distances and make a table out of these values. In 2016 we used this method to adjust the aim of our 2-axis turret based on how far from the goal we were.

Case Study: Aiming Using Vision

1. You can accurately and quickly aim your robot using only a limelight and your drivetrain.
2. We added the limelight to a robot, implemented the code and tuned it in less than 1hr.

Using high-framerate vision tracking, it is now possible to use the vision pipeline directly as the “sensor” in a PID control loop to guide your robot or turret. In order to test this idea we added a limelight to our 2017 FRC robot and made it aim at vision targets using nothing more than the drivetrain and the networks table data being reported by the limelight.

In this example, our test candidate was a 2017 FRC robot which uses a 6-wheel drivetrain with colson wheels. Here is a picture of us adding a limelight onto the robot in order to do this test.



Next we added some code to the robot which would run whenever the driver holds a button on the joystick. This robot used “tank” style driving so the `OperatorControl` function was generating a ‘`left_command`’ value and a ‘`right_command`’ value to control the left and right sides of the drivetrain. After the normal control code, we added a block of code like this:

```
float Kp = -0.1f; // Proportional control constant

std::shared_ptr<NetworkTable> table = NetworkTable::GetTable("limelight");
float tx = table->GetNumber("tx");

if (joystick->GetRawButton(9))
{
    float heading_error = tx;
    steering_adjust = Kp * tx;

    left_command+=steering_adjust;
    right_command-=steering_adjust;
}
```

Right off the bat, this mostly worked. The robot turns in the direction of the target automatically whenever you are holding the button. If you move the target around, the robot turns to follow the target. However, using the live video feed on the dashboard, we could see that there was one big problem: The robot wasn’t always making it all the way to perfectly align with the target. In some games with small targets, (like 2016 and 2017) this wouldn’t be good enough.

What we have implemented so far is a simple proportional control loop. We calculated the error in heading and multiplied that by a constant, thus making a motor command which is proportional to the error. As the error goes to

zero, our command will go to zero. The problem is that there is a lot of friction involved when the robot tries to turn. Very small commands will not turn the robot at all. At small angles, the command can become too small to actually move the robot. You might find that your robot reaches its target well when you start out with a large targeting error but it just can't aim at all if you start out really close.

There are a few ways to solve this problem but here is a really simple solution. We used a concept the "minimum command". If the error is bigger than some threshold, just add a constant to your motor command which roughly represents the minimum amount of power needed for the robot to actually move (you actually want to use a little bit less than this). The new code looks like this:

```
float Kp = -0.1f;
float min_command = 0.05f;

std::shared_ptr<NetworkTable> table = NetworkTable::GetTable("limelight");
float tx = table->GetNumber("tx");

if (joystick->GetRawButton(9))
{
    float heading_error = -tx;
    float steering_adjust = 0.0f;
    if (tx > 1.0)
    {
        steering_adjust = Kp*heading_error - min_command;
    }
    else if (tx < -1.0)
    {
        steering_adjust = Kp*heading_error + min_command;
    }
    left_command += steering_adjust;
    right_command -= steering_adjust;
}
```

Beware, if you set Kp or min_command too high, your robot can become unstable and can oscillate back and forth as it overshoots the target:

After some tuning on Kp and min_command should get your robot aiming directly at the target very accurately and quickly.

Case Study: Seeking

1. Use a seeking behavior to find a target that is not initially visible to the camera.

Implementing a seeking behavior can be useful in a few situations. In some games, the goal you are trying to score into may not be in a predictable location so you can't rely on dead-reconing autonomous code to get "close enough". In 2009 for example, the goals were attached to the opponents robots which also driving around. Another reason you might want to implement seeking is if your robot's drivetrain is not reliable/repeatable enough to put you in exactly the location you want. Sometimes swerve drives or omni-directional drives can have trouble driving to precise locations (this is not a knock on swerve or omni-directional drivetrains; it is just more challenging to get them to drive to repeatable locations). Some games have even allowed robot-robot interaction during autonomous mode. In 2006 we often engaged in autonomous duels where the defense robots would try to knock the offense robots out of scoring position. In this case you can't be sure of anything! Hopefully this gets you thinking about some ways that having your robot automatically seek for a goal that it can't initially see might be useful.

Luckily, adding seeking logic to your robot is very easy. To implement seeking, you simply have your robot scan for a target and once it finds it, start running the aiming code. We will use the 'tv' or 'target valid' value reported by limelight to know if we see a target or not.

```
std::shared_ptr<NetworkTable> table = NetworkTable::GetTable("limelight");
float tv = table->GetNumber("tv");
float tx = table->GetNumber("tx");

float steering_adjust = 0.0f;
if (tv == 0.0f)
{
    // We don't see the target, seek for the target by spinning in place at a
    ↪safe speed.
    steering_adjust = 0.3f;
}
else
{
    // We do see the target, execute aiming code
    float heading_error = tx;
    steering_adjust = Kp * tx;
}
```

(continues on next page)

(continued from previous page)

```
left_command+=steering_adjust;  
right_command-=steering_adjust;
```

Note that if your aiming code includes correcting the distance to the target, the robot behavior will be seeking followed by driving to the shooting distance as it aligns with the target. All of this can be achieved with the very simple code shown above.

CHAPTER 10

Case Study: Getting in Range

1. Have your robot automatically drive to a desired distance from the target.
2. A very simple trick using a calibrated cross-hair.

FRC robots often need to be positioned a specific distance from a scoring goal in order for their scoring mechanism to work well. You can use a simple proportional control loop along with limelight's calibrated cross-hair to very easily have your robot drive to a specific distance from the goal.

For this example, you should make sure that your robot is aimed at the goal. Later we will combine aiming and distance adjust into one function but for this example we are focusing only on the code that drives the robot to the correct distance.

Similar to the aiming example, here we show a function which can be run in your robot's update loop.

```
float KpDistance = -0.1f; // Proportional control constant for distance
float current_distance = Estimate_Distance(); // see the 'Case Study: Estimating_
↳Distance'

if (joystick->GetRawButton(9))
{
    float distance_error = desired_distance - current_distance;
    driving_adjust = KpDistance * distance_error;

    left_command += distance_adjust;
    right_command += distance_adjust;
}
```

With some tuning of the KpDistance setting, your robot should very quickly and accurately drive to the desired distance. Compared to aiming, driving to the correct distance is usually easier because most robots drive forward and backward much easier than they turn in place.

Next we'll describe a simple trick to make driving to the correct distance even easier. Instead of actually calculating the distance, you can use the limelight cross-hair. Just position your robot at your idea distance from the target and calibrate the y-position of the cross-hair. Now your y-angle will report as 0.0 when your robot is at the correct distance. Using this trick, you don't ever have to actually calculate the actual distance and your code can look something like this:

```
float KpDistance = -0.1f;

std::shared_ptr<NetworkTable> table = NetworkTable::GetTable("limelight");
float distance_error = table->GetNumber("ty");

if (joystick->GetRawButton(9))
{
    driving_adjust = KpDistance * distance_error;

    left_command += distance_adjust;
    right_command += distance_adjust;
}
```

If you need to change the shooting distance, just move your robot to the new distance and re-calibrate the limelight crosshair.

Here is an example of a robot using the above idea. Note how it backs up to the correct range automatically when the driver stops driving and presses the “aim” button:

Case Study: Aiming and Range at the same time.

1. Put aiming and range adjustment into one function.

This example uses code from the aiming and range adjustment examples and puts everything together into one simple function. Using this, you can get your robot “close” and then use code to automatically aim and drive to the correct distance.

```
float KpAim = -0.1f;
float KpDistance = -0.1f;
float min_aim_command = 0.05f;

std::shared_ptr<NetworkTable> table = NetworkTable::GetTable("limelight");
float tx = table->GetNumber("tx");
float ty = table->GetNumber("ty");

if (joystick->GetRawButton(9))
{
    float heading_error = -tx;
    float distance_error = -ty;
    float steering_adjust = 0.0f;

    if (tx > 1.0)
    {
        steering_adjust = KpAim*heading_error - min_aim_command;
    }
    else if (tx < 1.0)
    {
        steering_adjust = KpAim*heading_error + min_aim_command;
    }

    float distance_adjust = KpDistance * distance_error;

    left_command += steering_adjust + distance_adjust;
```

(continues on next page)

(continued from previous page)

```
right_command -= steering_adjust + distance_adjust;  
}
```

Case Study: 2017 Fuel Robot

1. Entire c++ robot program which implements 2017 fuel scoring.

In this example, we present a complete robot program which implements the 2017 boiler goal. We tested this program on the 987 steamworks robot and were able to easily score fuel in the boiler. For this example we aim and move into range using the robot's drivetrain. One thing that was immediately apparent to us is that limelight let us implement this feature in a tiny fraction of the amount of code and time compared to our real 2017 robot.

There are also a couple of other features in the code below such as the ability to blink the LEDs and dynamically toggle between multiple vision pipelines. This example also shows how to use a Talon SRX speed controller along with a magnetic encoder to control the RPM of the shooter wheel.

```
#include <iostream>
#include <string>

#include <Drive/DifferentialDrive.h>
#include <Joystick.h>
#include <SampleRobot.h>
#include <SmartDashboard/SendableChooser.h>
#include <SmartDashboard/SmartDashboard.h>
#include <Spark.h>
#include <Victor.h>
#include <Timer.h>
#include "ctre/phoenix/MotorControl/CAN/TalonSRX.h"

//
#define Vic_Drive_Left_1 0
#define Vic_Drive_Left_2 1
#define Vic_Drive_Right_1 2
#define Vic_Drive_Right_2 3

#define Tal_Shooter_Wheel 21
#define Tal_Shooter_Wheel_2 2
#define Tal_Intake 3
#define Tal_Uptake 9
```

(continues on next page)

(continued from previous page)

```

class Robot : public frc::SampleRobot
{
public:

Robot() :
    LeftDrive(Vic_Drive_Left_1),
    RightDrive(Vic_Drive_Right_1),
    LeftDrive2(Vic_Drive_Left_2),
    RightDrive2(Vic_Drive_Right_2),
    Shooter(Tal_Shooter_Wheel),
    Shooter_2(Tal_Shooter_Wheel_2),
    Intake(Tal_Intake),
    Uptake(Tal_Uptake)
{
    // This code sets up two Talons to work together and implement PID RPM control
    Shooter.SetControlMode(CTRE::MotorControl::ControlMode::kSpeed);
    Shooter_2.SetControlMode(CTRE::MotorControl::ControlMode::kFollower);
    Shooter_2.Set(Tal_Shooter_Wheel);
    Shooter.SetClosedLoopOutputDirection(false);
    Shooter.
↪SetStatusFrameRateMs(CTRE::MotorControl::CAN::TalonSRX::StatusFrameRateAnalogTempVbat,
↪10);

    Shooter.SetFeedbackDevice(CTRE::MotorControl::CAN::TalonSRX::CtreMagEncoder_
↪Relative);
    SetShooterConstants(.0002f, .00009f, .0000f, .000006f); // P,I,D,F constants_
↪for the shooter wheel
    Shooter.SetAllowableClosedLoopErr(10);
    Shooter.ConfigNominalOutputVoltage(+0.f,-0.f);
    Shooter.ConfigPeakOutputVoltage(+12.f,-12.f);
    Shooter.SelectProfileSlot(0);
    Shooter.SetSensorDirection(true);
    Shooter.SetVoltageRampRate(100);
    Shooter.SetIzone(300);
    Shooter_2.SetVoltageRampRate(100);
}

void RobotInit()
{
}

void Autonomous()
{
}

/*
 * Operator Control
 */
void OperatorControl() override
{
    while (IsOperatorControl() && IsEnabled())
    {
        // Read the joysticks!
        float left_command = m_stickLeft.GetY();
        float right_command = m_stickRight.GetY();

        // Get limelight table for reading tracking data
    }
}
    
```

(continues on next page)

(continued from previous page)

```

std::shared_ptr<NetworkTable> table = NetworkTable::GetTable(
↪ "limelight");

    float KpAim = 0.045;
    float KpDist = 0.0f; //0.09;
    float AimMinCmd = 0.095f;

    float targetX = table->GetNumber("tx", 0);
    float targetY = table->GetNumber("ty", 0);
    float targetA = table->GetNumber("ta", 0);

    // Aim error and distance error based on calibrated limelight cross-
↪ hair

    float aim_error = targetX;
    float dist_error = targetY;

    // Steering adjust with a 0.2 degree deadband (close enough at 0.2deg)
    float steeringAdjust = KpAim*aim_error;
    if (targetX > .2f) steeringAdjust += AimMinCmd;
    else if (targetX < -.2f) steeringAdjust -= AimMinCmd;

    // Distance adjust, drive to the correct distance from the goal
    float drivingAdjust = KpDist*dist_error;
    bool doTarget = false;

    if(m_stickLeft.GetRawButton(3))           // Aim using pipeline 0
    {
        doTarget = true;
        table->PutNumber("pipeline", 0);
    }
    else if (m_stickLeft.GetRawButton(2))     // Aim using pipeline 1
    {
        doTarget = true;
        table->PutNumber("pipeline", 1);
    }

    if(doTarget)                             // If auto-aiming, adjust_
↪ drive and steer
    {
        ShooterSetRPM(3300);
        left_command += drivingAdjust - steeringAdjust;
        right_command += drivingAdjust + steeringAdjust;
    }
    else
    {
        ShooterOff();
    }

    // Tank drive, send left and right drivetrain motor commands
    StandardTank(left_command, right_command);

    if(m_stickRight.GetRawButton(3))         // Suck in and shoot balls
    {
        IntakeIn();
        UptakeUp();
    }
    else if(m_stickRight.GetRawButton(2))     // Spit out balls

```

(continues on next page)

(continued from previous page)

```

        {
            IntakeIn();
            UptakeDown();
        }
        else // Leave the balls alone!
        {
            IntakeOff();
            UptakeOff();
        }
        if(m_stickLeft.GetRawButton(5)) // Joystick Button 5 = Flash
→the LEDs
        {
            table->PutNumber("ledMode", 2); //flash the lights
        }
        else
        {
            table->PutNumber("ledMode", 0); //turn the lights on
        }

        // wait for a motor update time
        frc::Wait(0.005);
    }
}

/*
 * Runs during test mode
 */
void Test() override {}

void StandardTank(float left, float right)
{
    LeftDrive.Set(-left);
    LeftDrive2.Set(-left);
    RightDrive.Set(right);
    RightDrive2.Set(right);
}

//
// Shooter Functions - uses talon PID to control shooter wheel RPM
// Set the P,I,D,F constants in the Talon, these values depend heavily on your
→mechanism
//
void SetShooterConstants(float p,float i,float d,float f)
{
    p *= 1024.f;
    i *= 1024.f;
    d *= 1024.f;
    f *= 1024.f;

    Shooter.SetP(p);
    Shooter.SetI(i);
    Shooter.SetD(d);
    Shooter.SetF(f);
}

//

```

(continues on next page)

(continued from previous page)

```

// Tell the talons our desired shooter wheel RPM
//
void ShooterSetRPM(float desrpm)
{
    Shooter.SetControlMode(CTRE::MotorControl::ControlMode::kSpeed);
    Shooter_2.SetControlMode(CTRE::MotorControl::ControlMode::kFollower);
    Shooter_2.Set(Tal_Shooter_Wheel);
    Shooter.Set(desrpm);
}

//
// Just set the power -1..+1, not currently using this now that RPM control is set up
//
void ShooterSetPower(float power)
{
    Shooter.SetControlMode(CTRE::MotorControl::ControlMode::kPercentVbus);
    Shooter_2.SetControlMode(CTRE::MotorControl::ControlMode::kPercentVbus);
    Shooter_2.Set(power);
    Shooter.Set(power);
}

//
// Turn off the shooter wheel
//
void ShooterOff()
{
    Shooter.SetControlMode(CTRE::MotorControl::ControlMode::kPercentVbus);
    Shooter_2.SetControlMode(CTRE::MotorControl::ControlMode::kPercentVbus);
    Shooter.Set(0.0f);
    Shooter_2.Set(0.0f);
}

//
// Intake Functions
//
void IntakeIn()
{
    Intake.Set(-.8f);
}
void IntakeOut()
{
    Intake.Set(.8f);
}
void IntakeShooting()
{
    Intake.Set(-1.0f);
}

void IntakeOff()
{
    Intake.Set(0);
}

//
// Uptake Functions
//
void UptakeUp()

```

(continues on next page)

(continued from previous page)

```
{
    Uptake.Set (-1.0f);
}
void UptakeDown()
{
    Uptake.Set (1.0f);
}
void UptakeOff()
{
    Uptake.Set (0);
}

private:

    // Robot drive system
    frc::Victor LeftDrive;
    frc::Victor RightDrive;
    frc::Victor LeftDrive2;
    frc::Victor RightDrive2;

    // shooter wheel
    CTRE::MotorControl::CAN::TalonSRX Shooter;
    CTRE::MotorControl::CAN::TalonSRX Shooter_2;
    CTRE::MotorControl::CAN::TalonSRX Intake;
    CTRE::MotorControl::CAN::TalonSRX Uptake;

    // Joystick inputs
    frc::Joystick m_stickLeft{0};
    frc::Joystick m_stickRight{1};
};

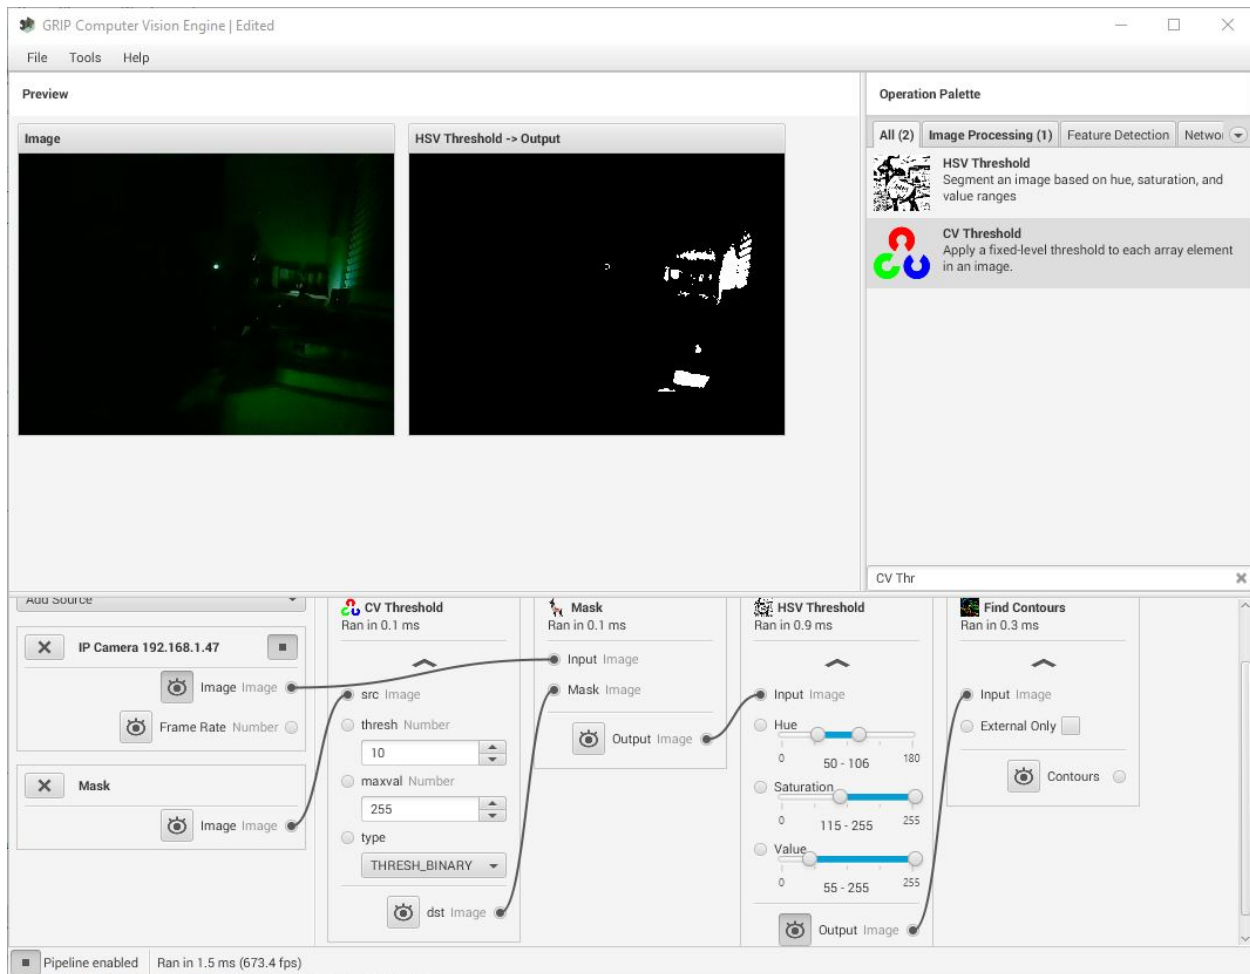
START_ROBOT_CLASS (Robot)
```

Using Grip to create Custom OpenCV Vision Pipelines

1. You can use GRIP to graphically design custom vision pipelines.
2. Your custom vision pipeline can be loaded into a limelight and run in hardware.

Grip is a program developed by WPI for developing vision pipelines. We have added a new export option to grip which allows you to export your custom vision pipeline and run it on a limelight camera with no extra programming needed. Carefully designed Grip pipelines can run at the full 90fps framerate of the limelight camera.

The first step in using GRIP is to download our build of Grip. It can be downloaded from our downloads page. When you run GRIP, you should see a screen like this:



You can learn more about GRIP on this web page: <https://wpilib.screenstepslive.com/s/currentCS/m/vision/463566-introduction-to-grip>

13.1 Accessing the Video Stream within GRIP

Limelight has an additional video stream on port 5802 which can be accessed primarily for use with GRIP or other applications like it. This video stream is uncompressed (or very lightly compressed) and has no cross-hair or other overlays drawn on the image. It is important to use this video stream for developing your GRIP pipeline because the normal Limelight video feed is optimized for very low bandwidth and latency which causes the image to be relatively low quality and can have significant compression artefacts. So to get your grip pipeline started, add an IP Camera input which uses the limelight as its input:

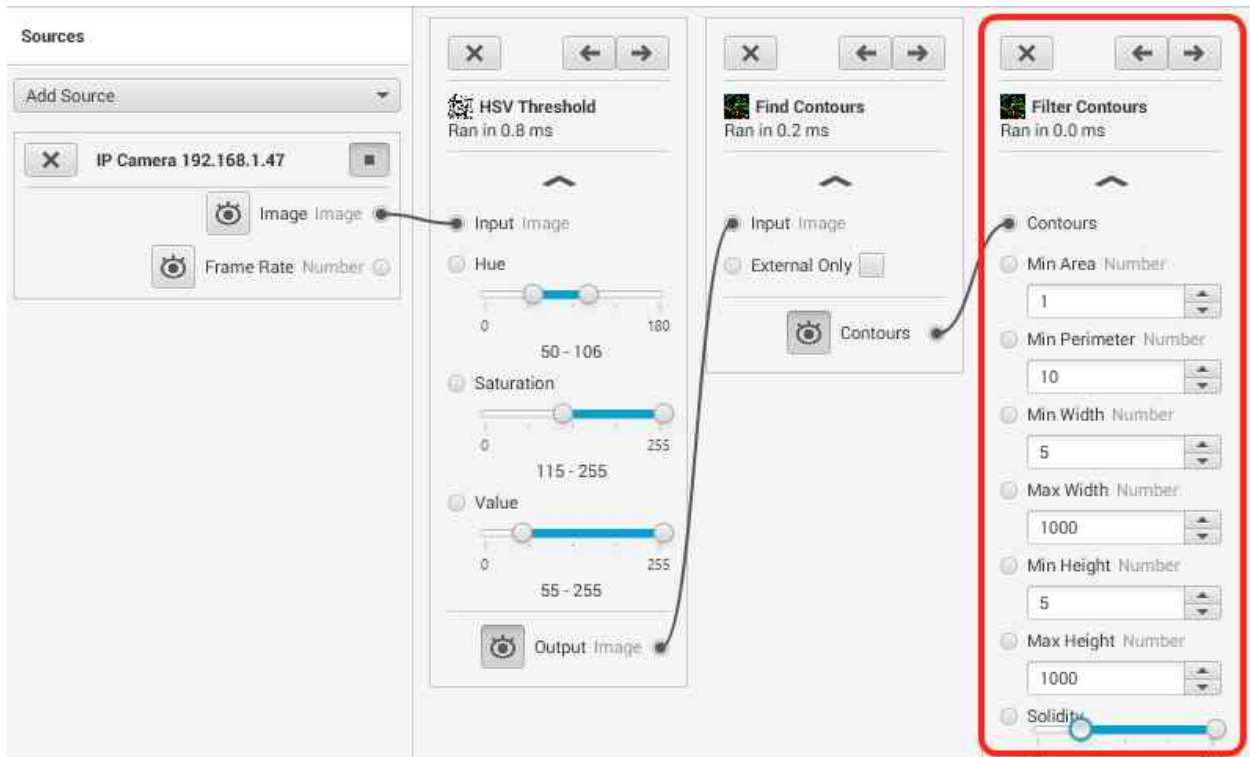


13.2 How Limelight Handles Inputs to the Pipeline

It is important to understand how your limelight camera interfaces with the inputs to your vision pipeline. Once your pipeline is loaded into the limelight camera, its inputs will be scanned. The first input is always assumed to be the video feed from the camera so it is important that you structure your pipeline in this way. Any additional inputs to the pipeline will receive upload buttons in the limelight UI which will allow you to upload additional images or data files as needed by your pipeline. For now it is important to realize that the video stream should be the first input in your Grip Pipeline.

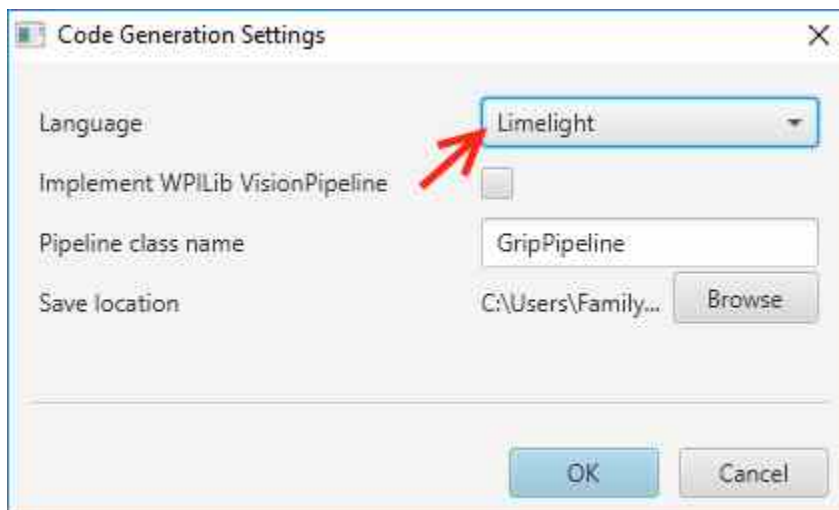
13.3 Generating output for the Limelight to use

When your grip pipeline is loaded onto the limelight, it will also be scanned for output data in the form of contours, blobs, or rectangles. If any of these outputs are found, the last one in the pipeline will be used as the 'results' of your vision pipeline. The outputs of your pipeline will be passed through the built-in contour filtering system on the limelight. If you want to do your own contour filtering within your GRIP pipeline, just set the limelight options to essentially accept any contour. In the example shown below, the output of the block highlighted in red would be passed on to the limelight targetting system. Notice that this example pipeline produces contours at an intermediate step and then filters them. Only the final set of contours (or blobs or rectangles) will be passed on to the limelight targetting system.



13.4 Code Generation for Limelight

When you are ready to try running your pipeline in hardware on the limelight, it is time to use the code-generation function of GRIP. We have added a new “Limelight” code generation target to the GRIP application. Using this option will produce a <pipeline>.ll file which can be loaded into the limelight camera.



Running GRIP Pipelines on Limelight

1. Upload your custom GRIP pipeline into a limelight pipeline.
2. Upload any additional inputs to your pipeline.
3. Observe the behavior of your GRIP pipeline.

Once you have used GRIP to design a custom vision pipeline and exported your pipeline to a .LL file, it is time to run it on the limelight camera hardware. The first step is to select which limelight pipeline you'd like to load your grip pipeline into. In order to do this, you use the UI at the top of the limelight interface to select a pipeline. Normally your robot (through network-tables) controls the currently active pipeline. Our first step is to disable that functionality so that we can select the pipeline to edit:



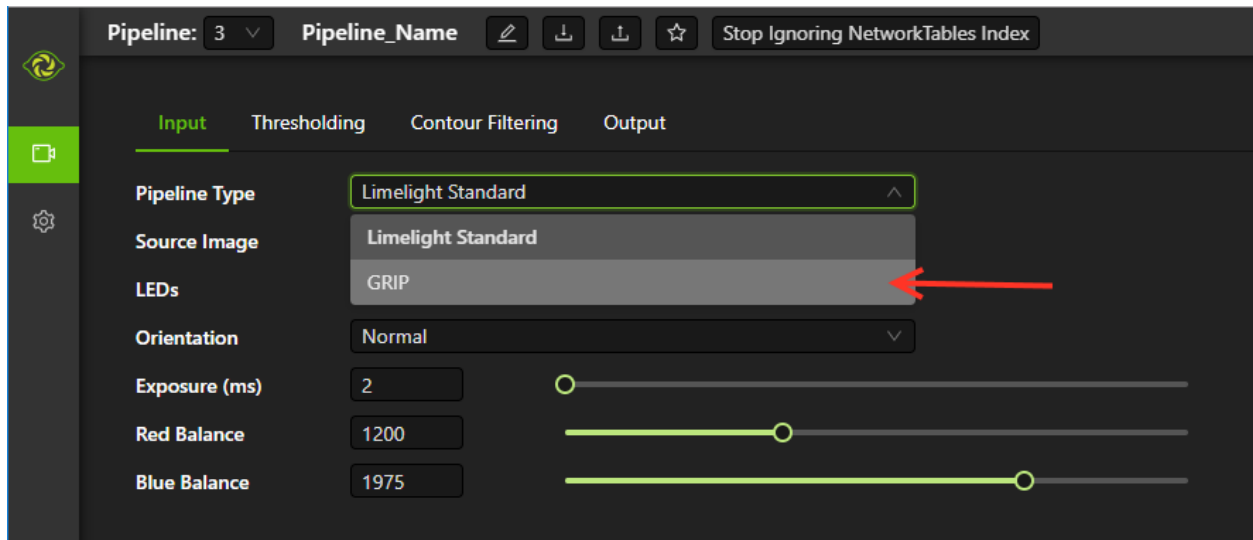
Next, change the currently active pipeline using the dropdown box highlighted below.



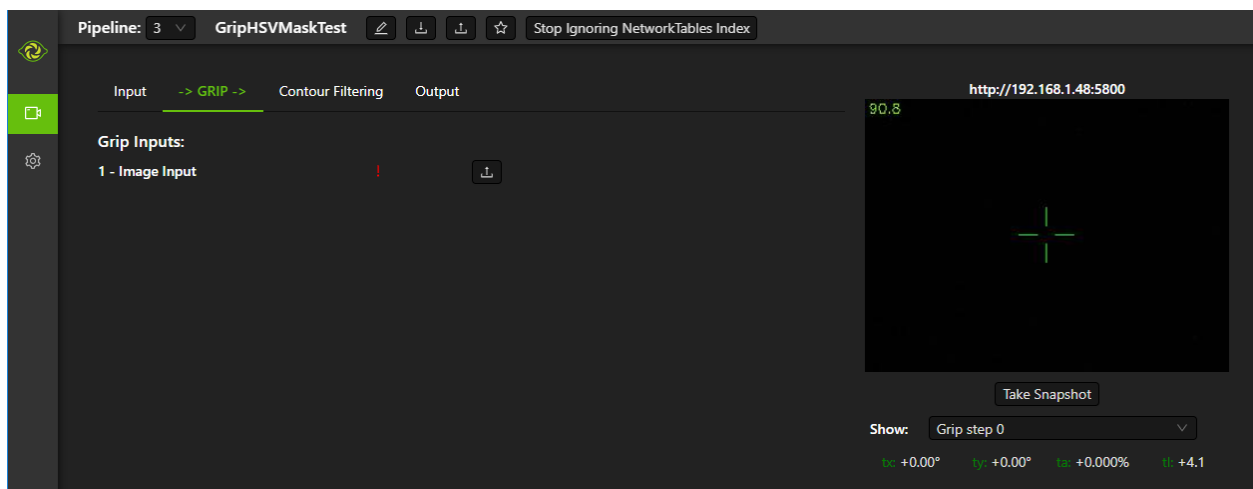
Next use the 'Upload' button to load your .LL file into this limelight pipeline.



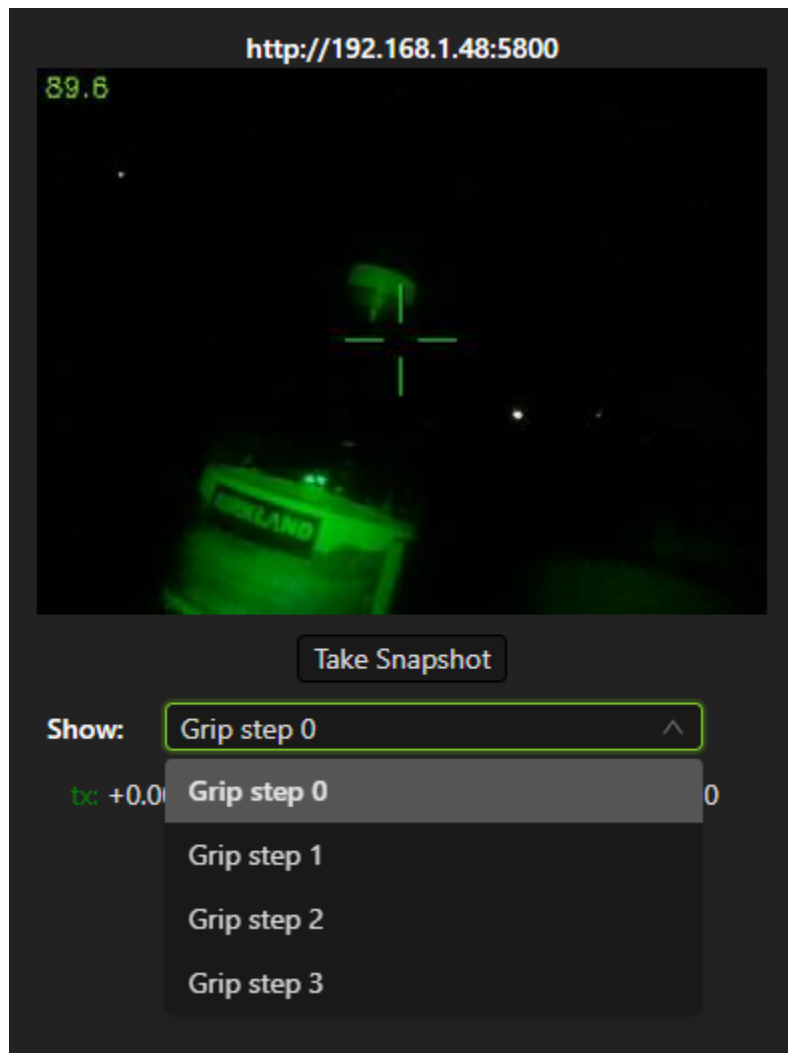
And finally, change the pipeline mode from “LimeLight Standard” to “GRIP”.



Once you have loaded a Grip pipeline, the UI changes slightly. In place of the 'Thresholding' tab, you will have a 'GRIP' tab. If your GRIP pipeline uses more than just the video feed as an input, you will have a list of additional inputs as shown below with upload buttons for uploading images or other data into them.



In addition, the video feed will have new options for displaying any intermediate images generated by your GRIP pipeline.



When you are using a GRIP pipeline, the settings on the Input, ContourFiltering, and Output tabs are still fully functional. If you would prefer to do your contour filtering within GRIP, you can simply “open up” all of the limelight contour filtering options to allow all contours to pass through.

Frequently Asked Questions

15.1 Why is limelight using a low (320x240) resolution?

Using a lower resolution allows us to run the vision pipeline at a high framerate. Limelight provides tracking data fast enough to use its outputs to directly drive a PID loop on your robot. This makes integrating vision into your robot extremely simple; limelight is just another sensor! Additionally, we have used even lower resolutions for the past two years in FRC. In our final event (China), our 2016 robot never mis-aimed a shot while using 160x120 resolution and aiming at one of the most difficult targets FRC has used to-date (the high goal in stronghold). *All* of its shots were auto-aimed; it was impossible to manually aim that robot.

15.2 What if the game calls for a different tracking algorithm?

We believe that Limelight will give you the tools to solve the vision tracking challenge for any FRC game. However if some challenge comes up which we have not anticipated, we will add the features needed for you to track the new type of target.

15.3 Why is there an extra usb port?

You can add a usb camera to limelight to act as a secondary driver or backup camera. The second camera stream will automatically appear alongside the primary camera stream. Limelight's video stream uses very little bandwidth so you can safely use this feature for a competitive advantage.

15.4 How do I view the video stream?

The video stream exists at <http://<Limelight's IP address or limelight.local>:5800>. Stream info is also posted to network tables so SmartDashboard and Shuffleboard (test LV dashboard) will automatically find it.

15.5 Are the eight LEDs bright enough?

The total light output of the LEDs is 400 lumens, and the LED cones actually increase the illuminance (functional light) of the device. To compare, the common two-ring setup has a total output of roughly 250 lumens.

15.6 Does Limelight support protocols other than NetworkTables?

Limelight also streams tx, ty, ta, and ts over the Serial (UART) pins at the back of the enclosure.

15.7 Does Limelight support Active PoE?

Limelight only supports passive PoE. Active POE requires a handshake between the sending and receiving POE devices. Limelight is designed to work with a passive POE injector. Just supply 12V by wiring the passive POE injector right into your PDP.

15.8 Should I plug Limelight into our VRM (Voltage Regulator Module)?

Limelight is designed to be connected to the main robot battery. It contains its own voltage regulators.

15.9 My robot has extreme voltage swings while driving, will this damage my Limelight?

We specifically designed Limelight to handle the “noisy” power available on a typical FRC robot. It contains the voltage regulators needed to generate the voltages it needs to operate as long its input voltage is above 6V.

15.10 Will Limelight’s LEDs dim when our robot’s voltage drops?

We designed the LED array on Limelight to produce a consistent brightness all the way down to 7.5V. So your tracking should remain consistent even when your robot motors are putting a heavy load on the battery. This can be an important feature for robots with high-traction drivetrains that aim by turning in place.

CHAPTER 16

Additional Resources

16.1 Team 254 Vision Talk

Team 254 gave an this excellent talk on how to do vision tracking using an android phone! Their talk is packed with very good information on the theory and implementation of vision tracking. <http://www.youtube.com/watch?v=rLwOkAJqlmo>

16.2 WPI Vision Processing Resources

These pages are loaded with information on various ways to implement your own vision processing pipeline. <https://wpilib.screenstepslive.com/s/currentCS/m/vision>

16.3 GRIP

Grip is a very popular tool that essentially lets you build an OpenCV vision pipeline using a drag-and-drop GUI. <https://wpilib.screenstepslive.com/s/currentCS/m/vision/l/463566-introduction-to-grip>

17.1 2019.2 (1/7/19)

2019.2 adds new features to better equip teams for the 2019 season.

- Intersection Filter
 -
 - The all-new intersection filter will allow teams to better distinguish between different groups of contours. The filter extends contours to infinity, and checks where they would intersect.
- Direction Filter
 -
 - The new direction filter will allow teams to better distinguish between contours of different orientations.
- Additional Changes and Fixes
 - IP Address is auto-posted to SmartDashboard/Shuffleboard for easier event troubleshooting. We have had almost no reports of this being an issue, but this serves as another contingency feature.
 - Pipeline name is auto-posted to SmartDashboard/Shuffleboard
 - Access the width and height of the rough bounding box in pixels with new networktables keys.
 - Access the longest side and shortest side of the perfect bounding box in pixels with new networktables keys.
 - “Sort Mode” now applies to dual-contours
 - “Sort Mode” is now fixed
 - 5802 stream is less taxing on the camera

17.2 2019.1 (12/18/18)

2019.1 fixes all major bugs reported during the beta

- Performance is back to a steady 90fps. A driver issue was the root cause of our performance problems during beta
- IP and hostname settings actually apply and “stick”
- Magic Wands have been re-ordered to match those of Limelight 2018 software
- We now support Grip inputs like masks through the web interface
- NetworkTables freeze fixed

17.3 2019.0 (12/10/18)

With 2019, we are introducing GRIP support, a brand-new interface, and a cleaner flashing procedure.

- Grip Support
 - Build pipelines with GRIP For Limelight, and export “LL Script” files to upload to your camera.
 - Masks and NetworkTables support will be added in a future update
 - Expose a bare video stream at <http://<limelighturl>:5802> for testing in grip
- All-new web interface
 - Smaller controls
 - More tooltips
 - Add the ability to turn off the LEDs from the interface
 - Move display combobox underneath stream for easier display switching
 - Faster communication to Limelight. New web tech allowed us to simplify other parts of our code.
- Flashing
 - We have migrated to “Balena Etcher”
 - Etcher is twice as fast and works on all platforms
 - Flash popups are fixed with the migration.
- Other
 - The LED Mode has been modified to allow for pipeline-specific LED Modes. LED MODE 0 is now “pipeline mode”, while mode 1 is “force off”
 - Faster boot times in automatic IP assignment mode.
 - Optimizations and minor bug fixes

17.4 2018.5 (3/28/18)

2018.5 fixes a critical issue that would prevent users from tuning pipelines on snapshots.

17.5 2018.4 (3/19/18)

2018.4 adds new contour sorting options. These are fairly important for cube tracking this year, as teams don't necessarily want to track the largest cube in view. In many cases, teams want to track the cube that is closest to their intakes. Many users have had to use the raw contours feature to implement their own sorting, so we want to make this as easy as possible.

- Contour Sort Mode
 - Select between “largest”, “smallest”, “highest”, “lowest”, “leftmost”, “rightmost”, and “closest” sort options.
 - We feel that many teams will make use of the “closest” option for cube tracking.
 -

17.6 2018.3 (2/28/18)

2018.3 fixes a major networktables reconnection bug which would cause NetworkTables settings changes to not propagate to Limelight. Thanks to Peter Johnson and the WPILib team for pinpointing and fixing the underlying NT bug. This was (as far as we know) the last high-priority bug facing Limelight.

Settings changes such as ledMode, pipeline, and camMode should always apply to Limelight. You should no longer need workarounds to change Limelight settings while debugging, after restarting robot code, and after rebooting the roborio.

17.6.1 Changes

- Fix major NT syncing issue which broke settings changes (ledMode, pipeline, and camMode) during LabView debugging, and after a reset/reboot of the roborio.
- Eye-dropper wand:
 - The eye dropper wand uses the same 10 unit window for Hue, but now uses a 30 unit window for saturation and value. This means that thresholding is more often a one-click operation, rather than a multi-step process.
- Snapshots
 - Setting the snapshot value to “1” will only take a single snapshot and reset the value to 0. Snapshotting is throttled to 2 snapshots per second.
 - Snapshot limit increased to 100 images.
 - Snapshot selector area is now scrollable to support 100 images.
 -

17.7 2018.2 (2/10/18)

2018.2 fixes all known streaming bugs with various FRC dashboards. It also makes Limelight easier to tune and more versatile during events.

17.7.1 Features

- Thresholding wands
 - Setup HSV threshold parameters in a matter of clicks
 - The “Set” wand centers HSV parameters around the selected pixel
 - The “Add” wand adjusts HSV parameters to include the selected pixel
 -
 - The “Subtract” wand adjusts HSV parameters to ignore the selected pixel
 -
- Snapshots
 -
 - Snapshots allow users to save what Limelight is seeing during matches or event calibration, and tune pipelines while away from the field.
 - Save a snapshot with the web interface, or by posting a “1” to the “snapshot” NetworkTables key
 - To view snapshots, change the “Image Source” combo box on the input tab. This will allow you to test your pipelines on snapshots rather than Limelight’s camera feed
 - Limelight will store up to 32 snapshots. It will automatically delete old snapshots if you exceed this limit.
- New Streaming options
 - We’ve introduced the “stream” NetworkTables key to control Limelight’s streaming mode. We’ve received requests for PiP (Picture-in-Picture) modes to better accommodate certain dashboards.
 - 0 - Standard - Side-by-side streams if a webcam is attached to Limelight
 - 1 - PiP Main - The secondary camera stream is placed in the lower-right corner of the primary camera stream.
 - 2 - PiP Secondary - The primary camera stream is placed in the lower-right corner of the secondary camera stream.
- Increase streaming framerate to 22fps
 - Look out for faster streams in an upcoming update
- Erosion and Dilation
 - Enable up to one iteration of both erosion and dilation.
 - Erosion will slightly erode the result of an HSV threshold. This is useful if many objects are passing through a tuned HSV threshold.
 - Dilation will slightly inflate the result of an HSV threshold. Use this to patch holes in thresholding results.
- Restart Button
 - Restart Limelight’s vision tracking from the web interface. This is only useful for teams that experience intermittent issues while debugging LabView code.

17.7.2 Optimizations

- Drop steady-state pipeline execution time to 3.5-4ms.

17.7.3 Bug Fixes

- Fix Shuffleboard streaming issues
- Fix LabView dashboard streaming issues

17.8 2018.1 (1/8/18)

- Red-Balance slider
- Blue-Balance slider
- Better default color balance settings
- Increased max exposure setting

17.9 2018.0 (1/3/18)

On top of a ton of new case studies, more detailed documentation, and a full example program for an autonomous STEAMWORKS shooter, the software has received a major upgrade.

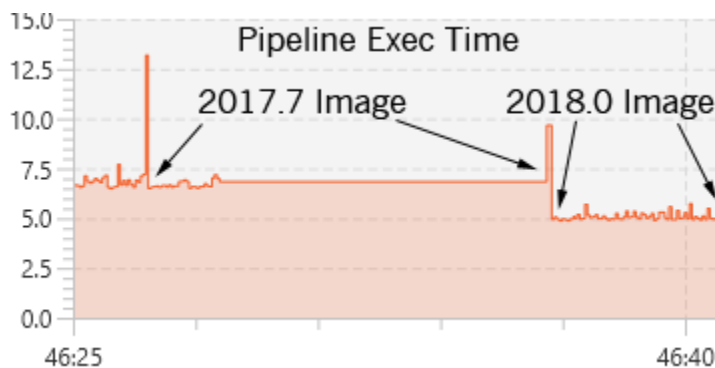
17.9.1 Features

- New Vision Pipeline interface:
 -
 - Add up to 10 unique vision pipelines, each with custom crosshairs, thresholding options, exposure, filtering options, etc.
 - Name each vision pipeline.
 - Mark any pipeline as the “default” pipeline.
 - Instantly switch between pipelines during a match with the new “pipeline” NetworkTables value. This is useful for games that have multiple vision targets (eg. the gear peg and boiler from 2017). This is also useful for teams that need to use slightly different crosshair options per robot, field, alliance, etc.
 - Download vision pipelines from Limelight to backup or share with other teams.
 - Upload vision pipelines to any “slot” to use downloaded pipelines.
- Target “Grouping” option:
 - Instantly prefer targets that consist of two shapes with the “dual” grouping mode”. “Single” and “Tri” options are also available
 -
- New Crosshair Calibration interface:
 - “Single” and “Dual” crosshair modes.
 - “Single” mode is what Limelight utilized prior to this update. Teams align their robots manually, and “calibrate” to re-zero targeting values about the crosshair.

- “Dual” mode is an advanced feature for robots that need a dynamic crosshair that automatically adjusts as a target’s area / distance to target changes. We’ve used this feature on some of our shooting robots, as some of them shot with a slight curve. This feature will also be useful for robots with uncentered and/or misaligned Limelight mounts.
- Separate X and Y calibration.
- Add Valid Target “tv” key to Network Tables.
- Add Targeting Latency “tl” key to Network Tables. “tl” measures the vision pipeline execution time. Add at least 11 ms for capture time.
- Draw additional rectangle to help explain aspect ratio calculation.
- Remove throttling feature, and lock Limelight to 90fps.
- Disable focusing on most web interface buttons. Fixes workflow problem reported by teams who would calibrate their crosshairs, then press “enter” to enable their robots.
- **Post three “raw” contours and both crosshairs to Network Tables.**
 - Access a raw contour with tx0, tx1, ta0, ta1, etc.
 - Access both raw crosshairs with cx0, cy0, cx1, cy1.
 - All x/y values are in normalized screen space (-1.0 to 1.0)
- Add “suffix” option to web interface. Allows users to add a suffix to their Limelights’ hostnames and Network-Tables (e.g. limelight-boiler). This feature should only be utilized if teams intend to use multiple Limelights on a single robot.
- Display image version on web interface

17.9.2 Optimizations

- Decrease networking-related latency to ~0.2 ms from ~10ms (Thanks Thad House)
- Move stream encoding and jpg compression to third core, eliminating 10ms hitch (25 - 30ms hitch with two cameras) seen every six frames.
- Drop steady-state pipeline execution time to 5ms with SIMD optimizations.



- New Latency testing shows 22 ms total latency from photons to targeting information.
- Upgrade Network Tables to v4 (Thanks Thad House)
- Optimize contour filtering step. Latency no longer spikes when many contours exist.
- Much improved hysteresis tuning.

- Significantly improve responsiveness of webinterface<->limelight actions.

17.9.3 Bugfixes

- Fix minor area value inaccuracy which prevented value from reaching 100% (maxed ~99%).
- Fix half-pixel offset in all targeting calculations
- Fix camera stream info not populating for NT servers started after Limelight's boot sequence. Regularly refresh camera stream info.
- Fix bug which caused aspect ratio to "flip" occasionally.
- Force standard stream output (rather than thresholded output) in driver mode.
- Fix bug which prevented LEDs from blinking after resetting Networking information

17.10 2017.7 (11/21/17)

- Improved contour sorting. Was favoring small contours over larger contours.
- New Coordinate system: Center is (0,0). ty increases as the target moves "up" the y-axis, and tx increases as the target moves "right" along the x-axis.
- More accurate angle calculations (Pinhole camera model).
- Display targeting info (tx, ty, ta, and ts) on webpage
- Default targeting values are zeros. This means zeros are returned if no target is in view.
- New side-by-side webpage layout. Still collapses to single column on small devices.
- Continuous slider updates don't hurt config panel performance.
- Aspect ratio slider scaled such that 1:1 is centered.

17.11 2017.6 (11/13/17)

- New Imaging tool. Tested on Win7, Win8 and Win10.
- Post camera stream to cameraserver streams. Works with smart dashboard camera streams, but shuffleboard has known bugs here
- Quartic scaling on area sliders, quadratic scaling on aspect ratio sliders. This makes tuning much easier
- Organize controls into "input", "threshold", "filter", and "output" tabs
- Continuous updates while dragging sliders
- Area sent to NT as a percentage (0-100)
- Image size down to 700MB from 2.1GB

17.12 2017.5 (11/9/17)

- Image size down to 2.1GB from 3.9GB
- Add driver mode and led mode apis
- Set ledMode to 0, 1, or 2 in the limelight table.
- Set camMode to 0 or 1 in the limelight table.
- Add ability to toggle between threshold image and raw image via web interface (will clean up in later release)
- Post camera stream to network tables under CameraPublishing/limelight/streams (will need a hotfix)
- Add skew to targeting information (“ts” in limelight table)
- Add base “CommInterface” in anticipation of more protocols

17.13 2017.4 (10/30/17)

- Lots of boot and shutdown bullet-proofing

17.14 2017.3 (10/25/17)

- Hue range is 0-179 from 0-255
- Decrease max log size, clear logs, clear apt cache

17.15 2017.2 (10/23/17)

- Manual ISO sensitivity
- Minimum exposure increased to 2

17.16 2017.1 (10/21/17)

- Read-only boot partition
- Mount all partitions with noatime option
- Disable all swap functionality (100 mb)
- Auto chmod 777 adminserver and visionserver from rc.local
- Smaller image
- “Convexity” changed to “Fullness”
- Exposure range set to 0-128 ms from 0-255 ms
- Support two cameras
- Fully support single-point calibration